

ĐẠI HỌC QUỐC GIA HÀ NỘI
TRƯỜNG ĐẠI HỌC CÔNG NGHỆ



Le Xuan An

Gradient-Routed Alignment with Capacity-aware Experts
for Multimodal Understanding

KHOÁ LUẬN TỐT NGHIỆP

Ngành: Khoa học máy tính

Giảng viên hướng dẫn: Tiến sĩ Nguyễn Văn Sơn

Giảng viên đồng hướng dẫn: Phó Giáo sư Võ Đình Hiếu

HÀ NỘI - 2026

**ĐẠI HỌC QUỐC GIA HÀ NỘI
TRƯỜNG ĐẠI HỌC CÔNG NGHỆ**



Le Xuan An

**Gradient-Routed Alignment with Capacity-aware Experts
for Multimodal Understanding**

KHOÁ LUẬN TỐT NGHIỆP

Ngành: Khoa học máy tính

Giảng viên hướng dẫn: Tiến sĩ Nguyễn Văn Sơn

Giảng viên đồng hướng dẫn: Phó Giáo sư Võ Đình Hiếu

HÀ NỘI - 2026

VIETNAM NATIONAL UNIVERSITY, HANOI
UNIVERSITY OF ENGINEERING AND TECHNOLOGY



Le Xuan An

**Gradient-Routed Alignment with Capacity-aware Experts
for Multimodal Understanding**

BACHELOR'S THESIS

Major: Computer Science

Advisor: Dr. Nguyen Van Son

Co-Supervisor: Associate Professor Vo Dinh Hieu

HANOI - 2026

Lời cảm ơn

Lời đầu tiên, tôi xin gửi lời cảm ơn chân thành và lòng biết ơn sâu sắc tới thầy TS. Nguyễn Văn Sơn và thầy PGS.TS. Võ Đình Hiếu vì đã hướng dẫn và hỗ trợ tôi về kiến thức chuyên môn vô cùng tận tình trong quá trình nghiên cứu và thực hiện khóa luận này.

Tôi cũng xin được gửi lời cảm ơn đến thầy cô, các anh chị và các bạn trong phòng thí nghiệm *iSE - Intelligent Software Engineering* đã luôn hỗ trợ tôi về kiến thức và chuyên môn trong quá trình làm khóa luận.

Cuối cùng, tôi xin gửi lời cảm ơn chân thành tới các thầy, cô trong trường Đại học Công Nghệ - Đại học Quốc gia Hà Nội đã tạo điều kiện cho tôi được học tập và nghiên cứu để có thể hoàn thành khóa luận với kết quả tốt nhất.

Assurance

I hereby affirm that this thesis is my own original work. I have not submitted it previously for any academic or professional award.

I completed this thesis through my own efforts, with guidance and support from Dr. Nguyen Van Son and Associate Professor Vo Dinh Hieu, in partial fulfillment of the graduation requirements. I confirm that all information, data and findings presented here are original and free from plagiarism.

I have acknowledged every source that I consulted and included each one in the references. I accept full responsibility for the accuracy and integrity of this thesis under the regulations of VNU University of Engineering and Technology.

Hanoi, Ngày 14 tháng 5 năm 2026

Student

Le Xuan An

Tóm tắt

Khóa luận này nghiên cứu bài toán chứng cất tri thức đa giáo viên không đồng nhất cho mô hình thị giác-ngôn ngữ (VLM). Mục tiêu là truyền tri thức từ một tập mô hình giáo viên đồng bằng sang một mô hình học sinh nhỏ hơn khi các giáo viên thuộc những họ mô hình khác nhau, sử dụng bộ tách từ khác nhau và có mức độ hữu ích thay đổi theo từng mẫu. Trong bối cảnh này, chứng cất trực tiếp theo chỉ số token không còn hợp lệ và phép lấy trung bình cố định giữa giáo viên có thể trộn tín hiệu hữu ích với tín hiệu không liên quan hoặc xung đột.

Để giải quyết bài toán trên, khóa luận đề xuất Gradient-Routed Alignment with Capacity-aware Experts (GRACE), một khung chứng cất cho truyền tri thức VLM đa giáo viên không đồng nhất. GRACE kết hợp ba cơ chế: chứng cất trie-Wasserstein trên cây byte UTF-8 dùng chung, định tuyến giáo viên thừa top- k từ trạng thái ẩn của mô hình học sinh và tinh chỉnh trọng số bằng độ đồng thuận gradient. Hàm mất mát trie-Wasserstein giúp so sánh phân phối đầu ra của giáo viên và học sinh giữa các bộ tách từ khác nhau. Định tuyến thừa chọn các giáo viên phù hợp với từng mẫu. Tinh chỉnh bằng độ đồng thuận gradient sau đó điều chỉnh trọng số đã định tuyến theo mức độ căn chỉnh giữa gradient chứng cất của từng giáo viên và gradient cross-entropy có giám sát. Thực nghiệm trên TextVQA, DocVQA và ChartQA đánh giá GRACE với hai mô hình học sinh SmolVLM-500M và SmolVLM-256M trong thiết lập bốn giáo viên không đồng nhất. Với SmolVLM-500M, GRACE đạt 69.38, 73.11 và 80.20 trên TextVQA, DocVQA và ChartQA, so với điểm gốc 56.75, 59.52 và 62.80 của mô hình học sinh. Các kết quả này tương ứng với mức cải thiện tương đối 22.26%, 22.83% và 27.71% so với mô hình gốc, đồng thời thu hẹp 55.03%, 39.53% và 55.34% khoảng cách ban đầu tới giáo viên có điểm cao nhất. Với SmolVLM-256M, GRACE đạt 60.73, 60.48 và 72.32 trên cùng ba bộ đánh giá. Trên các thực nghiệm được báo cáo, GRACE vượt qua tinh chỉnh trực tiếp, chứng cất logits cùng họ mô hình, ULD khác họ mô hình và các phương pháp đa giáo viên không đồng nhất được so sánh.

Từ khóa: mô hình thị giác-ngôn ngữ, chứng cất tri thức, chứng cất đa giáo viên, định tuyến giáo viên thừa, đồng thuận gradient, chứng cất Wasserstein khác bộ tách từ

Abstract

This thesis studies heterogeneous multi-teacher knowledge distillation for vision-language models (VLMs). The problem is to transfer supervision from a pool of frozen teacher models to a smaller student when the teachers come from different model families, use different tokenizers and vary in usefulness across samples. In this setting, direct token-index distillation is invalid and fixed teacher averaging can mix helpful teacher signals with irrelevant or conflicting ones.

To address this problem, the thesis proposes Gradient-Routed Alignment with Capacity-aware Experts (GRACE), a distillation framework for heterogeneous teacher-to-student VLM transfer. GRACE combines three mechanisms: trie-Wasserstein knowledge distillation over a shared UTF-8 byte trie, sparse top- k teacher routing from student hidden states and gradient-agreement refinement of routed teacher weights. The trie-Wasserstein loss makes teacher and student output distributions comparable across tokenizers. Sparse routing selects sample-relevant teachers. Gradient-agreement refinement then adjusts routed weights according to whether each teacher’s distillation gradient aligns with the supervised cross-entropy gradient.

Experiments on TextVQA, DocVQA and ChartQA evaluate GRACE with SmolVLM-500M and SmolVLM-256M students under a four-teacher heterogeneous pool. For SmolVLM-500M, GRACE reaches 69.38, 73.11 and 80.20 on TextVQA, DocVQA and ChartQA, compared with original student scores of 56.75, 59.52 and 62.80. These results correspond to relative gains of 22.26%, 22.83% and 27.71% over the original student and close 55.03%, 39.53% and 55.34% of the original gap to the highest-scoring teacher. For SmolVLM-256M, GRACE reaches 60.73, 60.48 and 72.32 on the same benchmarks. Across the reported experiments, GRACE improves over direct finetuning, same-family logits KD, cross-family ULD and the compared heterogeneous multi-teacher baselines.

Keywords: vision-language model, knowledge distillation, multi-teacher distillation, sparse teacher routing, gradient agreement, cross-tokenizer Wasserstein distillation

Table of Contents

Lời cảm ơn

Assurance	i
Tóm tắt	ii
Abstract	iii
Table of Contents	iv
List of Figures	vii
List of Tables	viii
List of Abbreviations	ix
Chapter 1 Introduction	1
Chapter 2 Background and Related Work	4
2.1 Vision-Language Models	4
2.1.1 Architecture Families	4
2.1.2 Autoregressive Answer Generation	5
2.2 Knowledge Distillation	6
2.2.1 Teacher-Student Objective	6
2.2.2 Temperature Scaling and Soft Targets	6
2.2.3 Answer-Token Supervision	7
2.2.4 Response-Based vs. Feature-Based Distillation	8
2.3 Tokenizer-Agnostic Distribution Matching	9
2.3.1 Wasserstein Distance	9
2.3.2 Tree-Structured Transport	10
2.4 Representation Matching	11
2.4.1 Hidden-State Alignment	11

2.4.2	Centered Kernel Alignment (CKA)	11
2.4.3	Limitations of Representation-Only Matching	12
2.5	Sparse Routing and Expert Selection	12
2.5.1	Mixture-of-Experts and Top- k Selection	12
2.5.2	Load Balancing Objectives	13
2.5.3	Sample-Adaptive Routing	14
2.5.4	Routing in Distillation Settings	14
2.6	Multi-Teacher Distillation	15
2.6.1	Weighted Teacher Aggregation	15
2.6.2	Gradient Conflict in Multi-Objective Optimization	16
2.7	Related Work	16
Chapter 3	Proposed Method	19
3.1	Method Overview and Problem Formulation	19
3.2	Trie-Wasserstein Cross-Tokenizer Distillation	23
3.2.1	Supervised Answer Tokens	23
3.2.2	Shared Byte Trie	23
3.2.3	Trie-Wasserstein Objective	25
3.2.4	Sparse Support and Residual Tail Edge	26
3.3	Teacher Routing Network	28
3.3.1	Router Architecture	28
3.3.2	Top- k Routing	29
3.3.3	Router Entropy and Z-Loss	30
3.4	Gradient-Agreement Refinement	30
3.4.1	Gradient Agreement Scoring	30
3.4.2	Agreement Scores and Active Set	31
3.4.3	Warmup, EMA Smoothing and Score Weights	31
3.4.4	Weight Blending and Fallback	32
3.5	Routed KD Aggregation and Final Objective	33
Chapter 4	Evaluation Methodology	35
4.1	Experimental Setup	35
4.2	Datasets	36
4.3	Training Configuration	37
4.4	Evaluation Metrics	38
4.5	Baselines and Compared Methods	38
4.6	Ablation Settings	40

Chapter 5 Experiment Results	42
5.1 Main Comparison Results	42
5.1.1 Baseline Results	42
5.1.2 Overall Comparison	43
5.2 Representation Similarity Analysis	51
5.3 Ablation Study	53
5.4 Efficiency Analysis	56
5.5 Discussion	58
Conclusion and Future Work	60
References	62
Appendix A Proof of the Closed-Form Trie-Wasserstein	69
Appendix B Training Configuration and Infrastructure	72
B.1 Baseline Implementation	73
B.2 Offline Teacher-Logit Preparation	74
B.3 Remote Logit Streaming During Training	76

List of Figures

2.1	Representative architecture families of modern vision-language models. . . .	5
2.2	Geometric behavior of Wasserstein distance.	10
2.3	Discrete optimal transport views.	11
3.1	Overview of the GRACE architecture.	20
3.2	Shared-trie tokenizer matching.	25
3.3	Sparse trie-Wasserstein support.	27
3.4	Teacher routing network.	28
3.5	Top- k routing dynamics.	29
3.6	Representative KD weighting regimes.	33
4.1	Representative dataset examples.	36
5.1	Grouped benchmark comparison.	45
5.2	DEX-AR overview.	46
5.3	Representative TextVQA attention example.	47
5.4	Uniform and GRACE training curves.	48
5.5	Routed-teacher load examples.	49
5.6	Final-layer CKA heatmaps.	52
5.7	Vision-layer CKA heatmaps.	53
B.1	R2 teacher-logit storage.	75
B.2	Training with remote teacher logits.	76

List of Tables

2.1	Summary of response-level and feature-level KD objectives.	9
3.1	Main notation.	22
4.1	Dataset split sizes.	37
4.2	Baseline groups.	39
5.1	Original SmolVLM student results.	42
5.2	Teacher and reference-model results.	43
5.3	Main benchmark comparison.	44
5.4	Relative gains of GRACE.	49
5.5	Gap-to-teacher analysis for SmolVLM-500M.	50
5.6	DocVQA component ablation.	54
B.1	Experiment parameters.	72
B.2	Logit-streaming settings.	76

List of Abbreviations

Abbr.	Full Term	Description
ANLS	Average Normalized Levenshtein Similarity	Evaluation metric used for DocVQA.
CE	Cross-Entropy	Supervised answer-token loss used during training.
CKA	Centered Kernel Alignment	Representation similarity measure used to compare hidden layers.
ChartQA	Chart Question Answering	Benchmark for question answering over charts.
DocVQA	Document Visual Question Answering	Benchmark for question answering on document images.
EOS	End of Sequence	Special token that marks the end of a generated answer span.
GRACE	Gradient-Routed Alignment with Capacity-aware Experts	Proposed multi-teacher distillation method in this thesis.
SHM-CKA	Hidden-State Matching with CKA	Baseline that combines logits distillation with hidden-state CKA matching.
JS	Jensen-Shannon	Symmetric divergence measure between probability distributions.
KD	Knowledge Distillation	Transfer of supervision from teacher models to a student model.
KL	Kullback-Leibler	Divergence used to compare probability distributions.
LLM	Large Language Model	Language model with large parameter scale and generation capability.
MLP	Multi-Layer Perceptron	Feedforward projection block used in the teacher routing network.
MoE	Mixture of Experts	Sparse routing framework that motivates teacher selection.
MTKD	Multi-Teacher Knowledge Distillation	Distillation setting that uses more than one teacher model.

Abbr.	Full Term	Description
OCR	Optical Character Recognition	Extraction of text from images or scanned documents.
OT	Optimal Transport	Distribution-matching view used to motivate trie-Wasserstein KD.
TWD	Tree-Wasserstein Distance	Tree-structured optimal-transport distance used as background for the proposed loss.
TextVQA	Text Visual Question Answering	Benchmark for question answering on natural images containing text.
ULD	Universal Logit Distillation	Tokenizer-agnostic distillation baseline based on ranked probability distributions.
VLM	Vision-Language Model	Model that processes visual and textual inputs together.
VLM-Eval-Kit	Vision-Language Model Evaluation Kit	Toolkit used to run standardized multimodal evaluation.
VQA	Visual Question Answering	Task of answering questions grounded in visual inputs.

Chapter 1

Introduction

Vision-language models (VLMs) have emerged as a **central direction** in modern AI because they integrate visual perception with language reasoning. By coupling powerful image encoders with autoregressive language decoders through learned multimodal connectors, these models can interpret visual content, read text embedded in images and reason across modalities in ways that neither vision systems nor language models alone can easily achieve [1–3]. This architectural synthesis has unlocked a **broad set of capabilities**, including extracting information from scanned documents, understanding visually rich scenes and grounding language instructions in physical environments [4–8].

Driven by neural scaling laws, the most capable VLMs are usually built on large language backbones trained with large multimodal corpora. This recipe reliably improves performance but also raises computational cost [9, 10]. As model size grows, the gap between what is achievable in research and what is deployable in practice widens accordingly. The **highest-scoring models are often the hardest to deploy** when latency, memory footprint, energy consumption and hardware constraints matter. In robotics and embodied AI, an agent must process visual observations, interpret language instructions and produce action decisions in real time, often on resource-constrained hardware [7, 8]. Similar constraints appear in mobile applications, document-processing pipelines and enterprise deployments, where users and operators need models with low latency, controlled cost and clear security boundaries. This widening **gap between frontier capability and deployment constraints** creates a central tension for applied VLM research.

Knowledge distillation offers a principled path through this tension by compressing the behavior of a large teacher model into a compact student. The student avoids training from scratch on the same large-scale corpora and learns from the teacher’s output distributions and internal representations, which carry richer information than hard labels alone [11, 12]. Distillation has been widely used to produce smaller models that **retain part of the capability of their larger counterparts** and recent lightweight VLMs demonstrate that this direction is viable for latency-sensitive and cost-sensitive deployment [13–15].

However, distillation in the VLM setting cannot be treated as a straightforward compression problem. Existing approaches each address part of the challenge, but **none covers it fully**. Response-based distillation aligns teacher and student output distributions effectively when both share a tokenizer, but **breaks down** when teacher and student use **different vocabularies** [11, 12]. Transport-based methods relax this assumption by comparing probability mass over decoded strings, but they are usually formulated for single-teacher transfer and do not determine which teacher should supervise a given sample [16–18]. Representation matching through hidden states or similarity measures such as CKA provides a tokenizer-independent signal, but does not directly optimize answer likelihood and **cannot resolve conflicts between teachers** [19]. Multi-teacher distillation methods show the value of pooling several teachers, but fixed or globally learned weights can mix useful supervision with low-utility or conflicting signals [20–22].

These limitations are most severe in text-rich visual reasoning tasks such as document question answering, chart interpretation and scene text understanding. Although recent open-source VLMs have narrowed the gap with proprietary systems on common benchmarks [23–25], smaller models continue to struggle when fine visual details, document structure and numerical relations must be captured. The central motivation of this thesis is to combine **tokenizer-agnostic** distribution matching, **sample-dependent** teacher selection and **gradient-aware conflict resolution** for heterogeneous VLM distillation.

This thesis evaluates **Gradient-Routed Alignment with Capacity-aware Experts (GRACE)** using VLM-Eval-Kit on TextVQA, DocVQA and ChartQA with SmolVLM-500M and SmolVLM-256M as student models. In the four-teacher setting, GRACE moves SmolVLM-500M closer to the highest-scoring teacher by closing 55.03%, 39.53% and 55.34% of original teacher-student gaps of 22.95, 34.38 and 31.44 points on TextVQA, DocVQA and ChartQA. This reduction supports deployment of smaller student VLMs in text-rich VQA scenarios. The study is guided by three research questions:

- **RQ1: Performance.** How does GRACE compare with direct fine-tuning, same-family logits KD, ULD and heterogeneous multi-teacher baselines on TextVQA, DocVQA and ChartQA benchmarks?
- **RQ2: Component Analysis.** How much do trie-Wasserstein KD, sparse top- k routing and gradient-agreement refinement contribute to the final performance?
- **RQ3: Efficiency.** What execution-time and computational overhead does GRACE introduce during training?

This thesis addresses this gap through three contributions to knowledge distillation for a pool of heterogeneous vision-language models.

First, trie-Wasserstein cross-tokenizer distillation loss. We formulate a tokenizer-agnostic KD objective that maps teacher and student token vocabularies onto a shared UTF-8 byte trie and computes a W_1 optimal-transport distance through subtree-mass imbalances. Compared with ranked-mass baselines such as ULD, the trie-Wasserstein loss preserves the lexical structure of decoded token strings: tokens that share a long byte prefix are cheaper to transport across than unrelated tokens. A sparse top- K_{TW} support with a residual tail edge keeps per-position cost bounded without discarding low-probability mass.

Second, gradient-agreement refinement of routed teacher weights. We introduce a post-routing correction mechanism that scores each routed teacher by comparing its KD gradient against the student’s supervised CE gradient with cosine similarity. The scores are stabilized with an exponential moving average, converted to weights through a score-softmax and geometrically blended with the router’s own weights. This refinement detects and suppresses teachers whose KD direction conflicts with the supervised objective, a failure mode that sparse routing alone cannot prevent.

Finally, GRACE as an integrated framework for heterogeneous multi-teacher VLM distillation. We combine the trie-Wasserstein loss, a sparse top- k teacher routing network and gradient-agreement refinement into a single end-to-end training objective. Empirical evaluation on TextVQA, DocVQA and ChartQA with SmolVLM-500M and SmolVLM-256M students under a four-teacher pool shows that GRACE improves over direct finetuning, same-family KD, cross-family ULD and the compared heterogeneous multi-teacher baselines, closing up to 55% of the teacher-student gap.

The remainder of this thesis is structured as follows. Chapter 2 provides background on VLMs, knowledge distillation, tokenizer-mismatched distribution matching, sparse routing and gradient conflict resolution. Chapter 3 presents GRACE in full, covering the trie-Wasserstein objective, sparse teacher router, gradient-agreement refinement and combined training objective. Chapter 4 describes the datasets, baselines, metrics and implementation details. Chapter 5 reports results across RQ1 to RQ3. The thesis concludes with the main findings and future directions.

Chapter 2

Background and Related Work

This chapter introduces the background concepts and related work used throughout the thesis. It first covers the building blocks of VLMs and knowledge distillation, then turns to the cross-tokenizer, routing and multi-teacher challenges. The final section connects these threads into the research gap addressed by this thesis.

2.1 Vision-Language Models

2.1.1 Architecture Families

Modern vision-language models often combine a vision encoder with a language model. On the visual side, many systems use a Vision Transformer pretrained with an image-text objective, commonly CLIP [26] or SigLIP [27], which converts an image into a sequence of patch embeddings that capture both local and global visual features. A connector module then fuses these visual tokens with text, using cross-attention, a Q-Former or a learned projector [1–3]. This design separates visual perception from language generation: the vision encoder produces a compact sequence of image features, while the language model provides the autoregressive prior used to answer questions. A useful way to distinguish VLM families is by where and how visual-language fusion happens, as shown in Figure 2.1. Dual-encoder models align image and text embeddings contrastively, cross-attention models let language tokens attend to visual features directly and projector-based models map visual embeddings into the language model hidden space and then decode with a standard causal language model. This thesis focuses on the projector-based generative design, because distillation targets output token distributions, not retrieval embeddings. In this setting, projected image embeddings are inserted into the language-model context as visual tokens, so answer generation remains ordinary next-token prediction conditioned on both the image representation and the textual question.

²Panels adapted from CLIP [26], Flamingo [1], BLIP-2 [2] and LLaVA [3].







Figure 2.1: Representative architecture families of modern vision-language models.²

2.1.2 Autoregressive Answer Generation

The distillation setting in this thesis assumes a generative VLM student that produces answer tokens autoregressively. For language modeling, the Transformer maps each hidden state $\mathbf{h}_j$ to next-token logits and probabilities [28]:

$$\mathbf{z}_j = \mathbf{W}_o \mathbf{h}_j + \mathbf{b}_o, \quad p_j(v) = \text{softmax}(\mathbf{z}_j)_v = \frac{\exp(z_{j,v})}{\sum_{u \in \mathcal{V}} \exp(z_{j,u})}. \quad (2.1)$$

At answer position j , the VLM can be abstracted as a mapping from the multimodal input and previous answer tokens to next-token logits:

$$\mathbf{z}_j = f_\theta(x, y_{<j}), \quad \mathbf{z}_j \in \mathbb{R}^{|\mathcal{V}|},$$

where x is a multimodal input and $\mathbf{z}_j$ denotes the logits used for autoregressive answer generation at position j . For an answer sequence $y = (y_1, \dots, y_L)$, the model factorizes the conditional probability as

$$p_\theta(y \mid x) = \prod_{j=1}^L p_\theta(y_j \mid x, y_{<j}). \quad (2.2)$$

Token-level distillation can then compare teacher and student distributions, such as $p_{t,j}$ and $p_{s,j}$, on the answer positions used for supervision [11, 28]. The loss is therefore naturally token-wise. At each answer position, the student distribution can be compared with either a hard label or a teacher distribution. This is important for VQA because short answers such as numbers, names and document fields may differ by only one or two tokens, so supervision must preserve position-level information, not only sequence-level accuracy after decoding the full answer. This output-centric view motivates the next sections: standard KD compares teacher and student answer-token distributions directly, while heterogeneous VLM distillation needs a tokenizer-agnostic matching objective.

2.2 Knowledge Distillation

2.2.1 Teacher-Student Objective

Knowledge distillation transfers information from a higher-capacity teacher model to a student model [11, 12]. Let $\mathcal{L}_{CE}$ denote the supervised cross-entropy loss of the student. A generic distillation objective takes the form

$$\mathcal{L} = \lambda_{CE}\mathcal{L}_{CE} + \lambda_{KD}\mathcal{L}_{KD}, \quad (2.3)$$

where $\mathcal{L}_{KD}$ is a teacher-guided loss. In token-level distillation, the teacher and student are compared on their output distributions, not only on hard labels. The student therefore learns richer distributional structure, the relative probabilities over incorrect classes often called *dark knowledge*, that hard labels discard [11]. For a single answer position, a common same-vocabulary form is

$$\mathcal{L}_{KD} = \tau^2 \text{KL}\left(\mathbf{p}_t^{(\tau)} \parallel \mathbf{p}_s^{(\tau)}\right), \quad (2.4)$$

where τ is the distillation temperature. The cross-entropy term anchors the student to the ground-truth answer, while the KD term transfers the teacher’s relative preference over alternative tokens. In the heterogeneous setting of this thesis, this template remains useful, but the comparison term must be changed because teachers and students may not share the same tokenizer.

2.2.2 Temperature Scaling and Soft Targets

Let $\mathbf{z}_s$ and $\mathbf{z}_t$ denote student and teacher logits. With temperatures τ_s and τ_t , the corresponding distributions are

$$\mathbf{p}_s = \text{softmax}\left(\frac{\mathbf{z}_s}{\tau_s}\right), \quad \mathbf{p}_t = \text{softmax}\left(\frac{\mathbf{z}_t}{\tau_t}\right). \quad (2.5)$$

Many classical KD methods compare $\mathbf{p}_s$ and $\mathbf{p}_t$ using divergences such as forward KL or Jensen-Shannon divergence. However, such approaches usually assume that teacher and student share the same vocabulary [11, 16, 29]. The temperature controls how much probability mass is exposed beyond the top answer token. When $\tau > 1$, the softmax distribution becomes flatter and low-rank alternatives receive non-negligible probability. When $\tau < 1$, the distribution becomes sharper and approaches hard-label supervision. This matters because the teacher may encode useful semantic alternatives in the tail of the distribution, for example related words, number formats or subword variants. The same softening idea is used later before applying tokenizer-independent distribution matching. In practice, the student and teacher temperatures may be chosen separately. Using τ_t for the teacher controls how much dark knowledge is exposed, whereas τ_s controls the smoothness of the student distribution being matched. For this reason, the temperatures are treated as part of the distillation interface throughout this thesis, not merely numerical constants inside the softmax.

2.2.3 Answer-Token Supervision

For multimodal question answering, distillation should focus on *answer tokens*, not *prompt tokens* [3]. The KD loss is therefore applied only on supervised answer positions after masking ignored tokens and truncating student and teacher answer spans to a common length. The result is a token-aligned distillation signal that remains compatible with different teacher and student tokenizers. Let Ω_b^{CE} denote the non-ignored student answer positions for sample b , so the supervised loss is

$$\mathcal{L}_{CE}^{(b)} = \frac{1}{|\Omega_b^{CE}|} \sum_{j \in \Omega_b^{CE}} -\log p_s(y_j \mid x_b, y_{<j}). \quad (2.6)$$

The same masking principle is used for KD positions. This prevents the distillation objective from spending capacity on prompt reconstruction, padding tokens or formatting artifacts and concentrates the teacher signal on the answer span that is scored by the evaluation benchmarks. For multiple teachers, the KD position set can be teacher-specific because teacher tokenizers and generated answer spans may differ. The training code therefore constructs $\Omega_{b,i}^{KD}$ separately for each teacher T_i , then averages only over valid aligned answer positions for that teacher.

2.2.4 Response-Based vs. Feature-Based Distillation

Knowledge distillation methods can be grouped by the type of teacher signal they transfer. Common response-level divergences and feature-level similarity losses used in knowledge distillation are collected in Table 2.1 [11, 29]. In the table, p and q denote teacher and student output distributions over token events t . $m = \frac{1}{2}(p + q)$ is the JS midpoint distribution. $\mathbf{u} = \Phi_T(f_T(x, y))$ and $\mathbf{v} = \Phi_S(f_S(x, y))$ are teacher and student feature vectors. The index r marks feature dimensions or normalized feature bins. *Response-based* distillation uses the teacher’s output distribution or generated response as supervision. In its standard form, the loss compares next-token distributions:

$$\mathcal{L}_{resp} = \frac{1}{|\Omega|} \sum_{j \in \Omega} \text{KL}(\mathbf{p}_{t,j}^{(\tau)} \parallel \mathbf{p}_{s,j}^{(\tau)}), \quad (2.7)$$

where Ω is the set of supervised answer positions. This objective is closely tied to the student’s final prediction layer and therefore provides a task-facing signal. It is the classical form of KD introduced by Hinton et al. [11] and remains common in language-model and VLM compression [12, 30].

Feature-based distillation aligns intermediate representations between student and teacher layers. Given student and teacher hidden states $\mathbf{H}_s^{(\ell)}$ and $\mathbf{H}_t^{(m)}$ at selected layers, a direct matching objective is

$$\mathcal{L}_{feat} = \left\| \phi_s(\mathbf{H}_s^{(\ell)}) - \phi_t(\mathbf{H}_t^{(m)}) \right\|_2^2, \quad (2.8)$$

where ϕ_s and ϕ_t are projection functions that make the hidden states comparable. Feature-based methods are useful when output vocabularies are incompatible, because the comparison can be defined in representation space without vocabulary-space matching [19, 31]. However, they are less directly connected to answer likelihood and may transfer teacher-specific internal structure that is not useful for the downstream task. This distinction motivates treating response-level and representation-level signals as complementary, not interchangeable.

Table 2.1: Summary of response-level and feature-level KD objectives.

Response-Level Divergences		Feature-Level Losses	
Type	$D(p, q)$	Function $\mathcal{L}_F$	Expression
Forward KLD	$\sum_t p(t) \log \frac{p(t)}{q(t)}$	L2-Norm Distance	$\ \mathbf{u} - \mathbf{v}\ _2$
Reverse KLD	$\sum_t q(t) \log \frac{q(t)}{p(t)}$	L1-Norm Distance	$\ \mathbf{u} - \mathbf{v}\ _1$
JS Divergence	$\frac{1}{2} \text{KL}(p\ m) + \frac{1}{2} \text{KL}(q\ m)$	Cross-Entropy Loss	$-\sum_r \mathbf{u}_r \log \mathbf{v}_r$

2.3 Tokenizer-Agnostic Distribution Matching

Same-vocabulary KD assumes that teacher and student probabilities are indexed by the same token set. Heterogeneous VLM distillation violates this: a teacher may tokenize the string “invoice” as one token while the student splits it into several subword pieces. Token-index KL then compares unrelated bins and is invalid [16].

The natural fix is to decode each token into a string and compare probability mass over decoded strings, not vocabulary indices. Let $\mathcal{V}_t$ and $\mathcal{V}_s$ be the teacher and student vocabularies with decoding maps $\psi_t : \mathcal{V}_t \rightarrow \mathcal{S}$ and $\psi_s : \mathcal{V}_s \rightarrow \mathcal{S}$ into a shared string space $\mathcal{S}$. Distribution matching is then defined over decoded strings, not vocabulary indices. This motivates transport-based objectives that replace same-vocabulary KL with a distance over probability mass on a shared lexical support.

2.3.1 Wasserstein Distance

Optimal transport compares two probability measures by the minimum cost required to move mass from one to the other [32]. For probability measures μ_s and μ_t , the p -Wasserstein cost is

$$W_p^p(\mu_s, \mu_t) = \min_{\gamma \in \Pi(\mu_s, \mu_t)} \int_{\Omega_s \times \Omega_t} \|x - y\|^p d\gamma(x, y), \quad (2.9)$$

where $\Pi(\mu_s, \mu_t)$ is the set of joint couplings with marginals μ_s and μ_t . For discrete weights $\mathbf{a}$, $\mathbf{b}$ and cost matrix $\mathbf{C}$, this becomes

$$W(\mathbf{a}, \mathbf{b}) = \min_{\mathbf{P} \in \mathcal{U}(\mathbf{a}, \mathbf{b})} \langle \mathbf{P}, \mathbf{C} \rangle, \quad \mathcal{U}(\mathbf{a}, \mathbf{b}) = \{\mathbf{P} \geq 0 \mid \mathbf{P}\mathbf{1} = \mathbf{a}, \mathbf{P}^\top \mathbf{1} = \mathbf{b}\}. \quad (2.10)$$

Unlike pointwise divergences, Wasserstein distances reflect displacement even when supports do not overlap, as Figure 2.2 illustrates.




Figure 2.2: Geometric behavior of Wasserstein distance.¹

2.3.2 Tree-Structured Transport

Dense OT requires a pairwise ground cost between all source and target bins, which is unavailable when vocabularies differ. A *tree metric* provides a structured alternative: probability mass is placed on a shared hierarchy and the distance is the sum of weighted *subtree-mass imbalances* [17, 32]. For a rooted tree $T = (V, E)$ with edge weights $w(e)$, one convenient form is

$$W_T(\mu, \nu) = \sum_{e \in E} w(e) |\mu(\Gamma(e)) - \nu(\Gamma(e))|, \quad (2.11)$$

where $\Gamma(e)$ denotes the subtree below edge e . The key point is that, on a tree metric, transport decomposes into edgewise subtree-mass imbalance and avoids a dense source-target coupling [17, 18]. The computation is efficient because the comparison follows the tree hierarchy directly and avoids materializing a full pairwise transport matrix. Figures 2.3a and 2.3b show the two commonly used OT views: a tree-metric view for fast subtree-mass comparisons and a matrix view for discrete transport plans. Tree-sliced OT extends this idea by averaging Wasserstein distances over tree metrics and shows that tree metrics preserve the geometric view of OT while admitting fast closed-form computations [18].

¹Adapted from Flamary and Courty [33].

¹Panel (a) is adapted from Takezawa et al. [17]. Panel (b) is adapted from Wikimedia Commons and Peyré and Cuturi [32, 34].




Figure 2.3: Discrete optimal transport views.¹

2.4 Representation Matching

2.4.1 Hidden-State Alignment

Representation matching transfers knowledge through intermediate hidden-state comparison, not final token distributions. Let $\mathbf{H}_s \in \mathbb{R}^{n_s \times d_s}$ and $\mathbf{H}_t \in \mathbb{R}^{n_t \times d_t}$ denote student and teacher hidden-state matrices. Because the sequence lengths and dimensions may differ, alignment usually requires pooling, projection, interpolation or layer selection:

$$\tilde{\mathbf{H}}_s = A_s(\mathbf{H}_s), \quad \tilde{\mathbf{H}}_t = A_t(\mathbf{H}_t), \quad \mathcal{L}_{hid} = \left\| \tilde{\mathbf{H}}_s - \tilde{\mathbf{H}}_t \right\|_F^2. \quad (2.12)$$

Here, A_s and A_t denote alignment operations that place student and teacher activations in a comparable space. This family of methods is naturally appealing for heterogeneous models because it avoids direct vocabulary comparison and can use information from image tokens, text tokens or fused multimodal states [31].

2.4.2 Centered Kernel Alignment (CKA)

Centered Kernel Alignment (CKA) compares representations through their similarity structure, not through elementwise hidden-state equality [19]. Given two activation matrices $\mathbf{X} \in \mathbb{R}^{n \times d_x}$ and $\mathbf{Y} \in \mathbb{R}^{n \times d_y}$ for the same n examples, the linear CKA score can be written as

$$\text{CKA}(\mathbf{X}, \mathbf{Y}) = \frac{\|\mathbf{Y}^\top \mathbf{X}\|_F^2}{\|\mathbf{X}^\top \mathbf{X}\|_F \|\mathbf{Y}^\top \mathbf{Y}\|_F}. \quad (2.13)$$

Higher CKA values indicate that the two representations induce similar geometry over the examples. Because CKA is invariant to isotropic scaling and orthogonal transformations, it is often more stable than direct hidden-state matching when the teacher and student have different hidden dimensions. In this thesis, CKA is therefore relevant both as a possible representation-level distillation signal and as an analysis tool for measuring how closely student features follow teacher features.

2.4.3 Limitations of Representation-Only Matching

Representation matching does not require a shared tokenizer, but it also does not directly optimize answer probabilities. A hidden-state objective can make the student resemble the teacher internally while still leaving the final answer distribution poorly calibrated. In a multi-teacher setting, the problem is sharper: two teachers may have useful but incompatible internal representations and a representation-only objective does not specify which teacher should dominate for a given input. It also does not detect whether a teacher’s representation loss produces a gradient that agrees with the supervised answer loss. For these reasons, representation matching is better viewed as a complementary signal or diagnostic baseline, not a complete solution to heterogeneous multi-teacher VLM distillation.

2.5 Sparse Routing and Expert Selection

2.5.1 Mixture-of-Experts and Top- k Selection

Recent sparse expert models route each sample or token through only a subset of available experts, improving efficiency and allowing specialized computation [35–37]. Let $\mathbf{h}^{(b)}$ denote a pooled representation for sample b . A router network $G_\phi(\cdot)$ produces gate logits

$$\mathbf{r}^{(b)} = G_\phi(\mathbf{h}^{(b)}), \quad \tilde{\mathbf{r}}^{(b)} = \text{AdjustedScores}(\mathbf{r}^{(b)}), \quad \mathbf{p}^{(b)} = \text{softmax}(\tilde{\mathbf{r}}^{(b)}), \quad (2.14)$$

where $p_i^{(b)}$ is the initial routing weight for expert i . The adjusted scores $\tilde{\mathbf{r}}^{(b)}$ may include expert bias, training noise or temperature scaling before the routing softmax. *Top- k routing* retains only the highest-scoring experts and zeros out all others. *Capacity constraints* prevent one expert from receiving too many assignments. If $\mathcal{S}_b = \text{TopK}(\mathbf{p}^{(b)}, k)$ is the selected expert set, the sparse expert output can be written as

$$\mathbf{y}^{(b)} = \sum_{i \in \mathcal{S}_b} \alpha_i^{(b)} E_i(\mathbf{h}^{(b)}), \quad \alpha_i^{(b)} = \frac{p_i^{(b)}}{\sum_{j \in \mathcal{S}_b} p_j^{(b)}}. \quad (2.15)$$

Only the selected experts are evaluated, so the activated computation can remain small even when the total number of expert parameters is large. The same abstraction is useful for teacher routing: each teacher can be treated as an expert that contributes a weighted distillation signal for a subset of samples.

2.5.2 Load Balancing Objectives

Load balancing prevents routing collapse onto a small subset of experts. In a sparse MoE layer, collapse is not only a modeling problem but also a systems problem: the overloaded experts determine the step latency, while under-used experts receive few gradients and remain poorly specialized [35, 38, 39]. Let a minibatch contain T routed items and let N denote the number of experts. For item b , the router produces probabilities $\mathbf{p}^{(b)} \in \Delta^{N-1}$. A top- k router induces the binary assignment mask

$$\mathbf{a}_i^{(b)} = \mathbf{1}\{i \in \text{TopK}(\mathbf{p}^{(b)}, k)\}, \quad n_i = \sum_{b=1}^T \mathbf{a}_i^{(b)}, \quad (2.16)$$

where n_i is the number of routed items assigned to expert i . Balanced routing corresponds to $n_i \approx kT/N$. When n_i is much larger than this value for a few experts, the router has effectively reduced the model to a small ensemble of busy experts while wasting the capacity of the remaining ones. The routing statistics used by many auxiliary losses can be written as a *soft* load and a *hard* load:

$$\bar{p}_i = \frac{1}{T} \sum_{b=1}^T p_i^{(b)}, \quad \bar{q}_i = \frac{1}{kT} \sum_{b=1}^T \mathbf{a}_i^{(b)}. \quad (2.17)$$

Here, $\bar{p}_i$ measures how much probability mass the router gives to expert i before hard selection, whereas $\bar{q}_i$ measures how often the expert is actually selected after sparse routing. The Switch Transformer balancing loss uses their scaled dot product [36]:

$$\mathcal{L}_{bal} = \alpha_{bal} N \sum_{i=1}^N \bar{p}_i \bar{q}_i. \quad (2.18)$$

The factor N keeps the magnitude approximately independent of the number of experts because the uniform solution gives $\sum_i (1/N)(1/N) = 1/N$. The hard load $\bar{q}_i$ is not differentiable through the top- k decision, but $\bar{p}_i$ is differentiable. The loss therefore gives the router a gradient signal that discourages assigning high average probability to already busy experts. In this thesis, the GRACE router uses entropy and z -loss terms in place of a hard capacity constraint, so load balancing is understood here as a stabilizer for sparse teacher selection, not a replacement for the task or distillation loss.

2.5.3 Sample-Adaptive Routing

Sparse routing is useful because the highest-utility expert can depend on the input. In a generic sample-adaptive router, the selected expert set is a function of a sample representation:

$$\mathcal{S}_b = \text{TopK}(\text{softmax}(G_\phi(\mathbf{h}^{(b)})), k). \quad (2.19)$$

The output or loss is then formed only from experts in $\mathcal{S}_b$. This differs from fixed ensembling, where every expert contributes to every example with the same global weight. Sample-adaptive routing is especially relevant when inputs have different structure. In multimodal QA, one sample may require OCR and document layout reasoning, another may require chart interpretation and another may depend mainly on natural image recognition. A router can learn to give different experts or teachers different responsibility across these cases. The mathematical distinction can be stated compactly. A global weighting scheme uses one vector $\boldsymbol{\pi} \in \Delta^{N-1}$ for the full training set:

$$\mathcal{L}^{(b)} = \sum_{i=1}^N \pi_i \mathcal{L}_i^{(b)}. \quad (2.20)$$

A sample-adaptive scheme uses $w_i^{(b)}$:

$$\mathcal{L}^{(b)} = \sum_{i=1}^N w_i^{(b)} \mathcal{L}_i^{(b)}, \quad w_i^{(b)} = 0 \text{ for } i \notin \mathcal{S}_b. \quad (2.21)$$

The sparsity constraint limits the number of active experts and makes the teacher-selection decision interpretable. It also reduces the chance that low-utility or irrelevant teachers dominate the update for a sample.

2.5.4 Routing in Distillation Settings

In distillation, frozen teachers can play the role of experts. The router does not combine hidden activations. Its role is to combine teacher-guided losses. If T_i is teacher i , then teacher T_i contributes a loss $\mathcal{L}_{KD}^{(b,i)}$ for sample b . A routed distillation loss takes the form

$$\mathcal{L}_{route}^{(b)} = \sum_{i \in \mathcal{S}_b} \omega_i^{(b)} \mathcal{L}_{KD}^{(b,i)}. \quad (2.22)$$

This view connects MoE selection to multi-teacher KD: teacher routing is a conditional computation mechanism over supervision signals. Unlike standard MoE, the experts here are frozen, so only the student and router parameters receive gradient updates. Routing in a distillation setting has two additional complications. First, teacher quality is not a fixed property of the teacher alone. It depends on the input, the current student state and the task objective. Second, a high router score does not guarantee that the selected teacher produces a gradient aligned with supervised learning. This means routing is useful but not a complete solution to teacher selection by itself.

2.6 Multi-Teacher Distillation

2.6.1 Weighted Teacher Aggregation

Let $\mathcal{T} = \{T_1, \dots, T_N\}$ be a pool of frozen teachers. For sample b , teacher T_i produces a per-teacher distillation loss $\mathcal{L}_{KD}^{(b,i)}$. In principle, one could average these losses uniformly,

$$\mathcal{L}_{KD}^{(b)} = \frac{1}{N} \sum_{i=1}^N \mathcal{L}_{KD}^{(b,i)}, \quad (2.23)$$

but this assumes every teacher is equally useful for every sample. In practice, different teachers may provide conflicting guidance because their predictions, specialized strengths and KD gradients need not agree on the same training sample. Some teachers may help the supervised objective on a given sample, while others may push the student in a different direction [40]. The sparse-routing framework introduced in Section 2.5 applies directly here by treating teachers as experts whose usefulness varies across samples [20–22]. More generally, an adaptive multi-teacher objective can be written as

$$\mathcal{L}_{KD}^{(b)} = \sum_{i=1}^N w_i^{(b)} \mathcal{L}_{KD}^{(b,i)}, \quad w_i^{(b)} \geq 0, \quad \sum_{i=1}^N w_i^{(b)} = 1. \quad (2.24)$$

Uniform averaging is the special case $w_i^{(b)} = 1/N$. A routed method lets the weights depend on the sample, the student state and the current teacher-selection policy. This is useful when teachers specialize differently: one teacher may score higher on documents, another on charts and another on natural scene text. The main risk is that adaptive weighting can collapse to one teacher.

2.6.2 Gradient Conflict in Multi-Objective Optimization

Even after sparse routing selects a teacher subset, the remaining teachers may still induce KD gradients that disagree with the supervised objective. To see why, consider the standard combined distillation loss $\mathcal{L} = \mathcal{L}_{CE} + \alpha \mathcal{L}_{KD}$. This objective minimizes a supervised term and a teacher-guided term together. Because both terms share the same student parameters, the combined update direction depends on whether the gradients $\nabla_{\theta}\mathcal{L}_{CE}$ and $\nabla_{\theta}\mathcal{L}_{KD}$ agree or conflict. When these two gradient directions are aligned, the combined update reinforces both objectives. When they oppose each other, the resulting update is a compromise that may degrade one or both objectives. This is an instance of *multi-objective optimization*: a single parameter update must serve competing objectives simultaneously [40, 41]. In multi-teacher distillation, the problem is compounded because each teacher contributes its own KD gradient. A routing mechanism can select a small teacher subset, but selection alone cannot guarantee that the remaining KD gradients support the supervised task update.

2.7 Related Work

Previous studies can be organized into four methodological groups, each addressing one part of heterogeneous VLM distillation.

Output-level KD and cross-tokenizer matching. Response-based KD transfers teacher output distributions to the student and has been used for language-model and multimodal compression [11, 12, 30, 42, 43]. Its main strength is that the supervision is task-facing: the loss directly acts on the answer-token distribution that determines generated answers. The limitation is that standard KL or JS losses usually assume a shared vocabulary and matched token indices. Tokenizer-agnostic and optimal-transport methods reduce this restriction by comparing probability mass over decoded strings or structured supports [16–18, 31]. These methods make cross-tokenizer transfer possible, but they are mainly formulated for single-teacher transfer and do not determine how to select or weight multiple heterogeneous teachers.

Representation matching and CKA-based alignment. Representation-level methods avoid output-vocabulary mismatch by aligning hidden states, attention maps or similarity structures [19, 31]. Their strength is tokenizer independence: teacher and student can have different vocabularies while the comparison is still defined in hidden-state space. CKA also gives a useful diagnostic view of whether two models organize examples similarly. The limitation is that representation similarity is an indirect proxy for task behavior. High hidden-state similarity does not guarantee matching answer probabilities after the language head and representation losses do not decide which teacher should dominate for a given sample.

Multi-teacher KD and sparse teacher routing. Multi-teacher KD studies how several teachers can supervise one student [20–22]. Its strength is teacher diversity: different teachers can provide different signals for documents, charts or scene text. Sparse expert routing adds a sample-dependent selection mechanism by choosing only a small active subset [35, 36, 39]. This improves scalability because training does not need to use every teacher for every sample. The limitation is that routing estimates relevance, not objective compatibility. A routed teacher can still produce a KD gradient that conflicts with the supervised answer loss and fixed averaging can dilute useful supervision with lower-quality signals.

Gradient-conflict and compatibility methods. Gradient-based filtering methods study whether teacher or auxiliary objectives agree with the supervised update direction [40, 41]. Their strength is that they measure objective compatibility directly in parameter-gradient space, which is closer to the training update than static teacher scores. The limitation is scope: these methods are not designed as cross-tokenizer multi-teacher VLM distillation objectives and they do not define a tokenizer-agnostic output loss for heterogeneous teacher logits.

Taken together, these groups leave four limitations for the setting studied in this thesis. First, same-vocabulary KD is not valid when teacher and student tokenizers differ. Second, cross-tokenizer methods do not by themselves solve teacher selection or multi-teacher weighting. Third, representation matching can be tokenizer-independent but may have lower task accuracy when representation similarity does not translate into answer-token agreement. Fourth, multi-teacher averaging and sparse routing improve scalability but can still use teacher signals whose gradients conflict with the supervised objective. These limitations can appear as invalid loss definitions under tokenizer mismatch, high transport or teacher-inference cost without sparsification, lower task accuracy when representation similarity fails to match answer behavior and limited scalability when every teacher is used for every sample. Addressing these limitations together requires a method that combines tokenizer-agnostic output matching, sample-adaptive teacher selection and objective-aware post-routing weighting.

Chapter 3

Proposed Method

This chapter presents **Gradient-Routed Alignment with Capacity-aware Experts (GRACE)** by first defining the heterogeneous multi-teacher distillation setting, then introducing the **trie-Wasserstein KD loss**, **teacher routing network**, **gradient-agreement refinement mechanism** and **final training objective**. The method targets vocabulary mismatch, variable teacher usefulness across samples and KD gradients that conflict with the supervised objective, with each component correcting one failure that the previous component cannot fix.

3.1 Method Overview and Problem Formulation

We consider multimodal understanding tasks that require reading image text, understanding the question and generating an answer sequence [4–6]. The training set is $\mathcal{D} = \{(x_n, y_n)\}_{n=1}^{N_{\text{data}}}$, where x_n is a multimodal input and y_n is the target answer sequence. Let S_θ denote the student VLM with trainable parameters θ , let G_ϕ denote the trainable teacher router with parameters ϕ and let $\mathcal{T} = \{T_i\}_{i=1}^N$ be the frozen teacher pool. For each sample b , the student produces logits $\mathbf{z}_s^{(b)}$ and teacher T_i provides teacher logits $\mathbf{z}_t^{(b,i)}$. This formulation treats teacher choice as part of the learning problem, not as a fixed preprocessing decision. In a heterogeneous teacher pool, a model that is useful for one image-question pair may be less useful for another. A teacher specialized for document layout may provide useful supervision for forms and receipts, while a different teacher may be more suitable for natural scene text or chart reasoning. At the same time, teacher logits cannot be compared to student logits by token index when their tokenizers differ. The method must therefore solve three coupled problems during training: compare teacher and student output distributions across different vocabularies, select a small set of relevant teachers for each sample and suppress teacher signals whose gradients conflict with the supervised answer objective. GRACE addresses these problems by combining tokenizer-agnostic trie-Wasserstein matching, student-conditioned sparse teacher routing and gradient-agreement refinement.




Figure 3.1: Overview of the GRACE architecture.

Figure 3.1 summarizes the full training pipeline. The diagram should be read from left to right. The left block contains the student, the frozen teacher pool and the sample-level quantities produced by the forward pass: student logits, hidden states and the supervised cross-entropy loss. The middle block contains the two mechanisms that produce candidate teacher supervision. Trie-Wasserstein KD computes a per-teacher distillation loss over a shared byte trie, enabling tokenizer-agnostic output matching even when teacher and student vocabularies are different. In parallel, the teacher routing network maps pooled student hidden states to sparse routing weights and selects a small subset of candidate teachers for the current sample. The right block applies gradient-agreement refinement. It compares each routed teacher’s KD gradient with the supervised CE gradient, smooths the agreement signal and converts the routed scores into final KD weights.

The method is organized around four components. **Trie-Wasserstein KD** computes per-teacher distillation losses over a shared byte trie, so tokens are matched by decoded lexical structure, not by vocabulary index. The **teacher routing network** selects a sparse subset of candidate teachers from student-side hidden states at each training step. **Gradient-agreement refinement** adjusts the routed teacher weights based on how well each teacher’s distillation gradient aligns with the supervised objective. Finally, the **training objective** combines supervised cross-entropy, routed KD and router auxiliary losses into a single end-to-end update. Together, they extend sparse expert routing and token-level logit supervision to the heterogeneous multi-teacher VLM setting, where teachers and student may use different tokenizers and vocabularies [35, 36, 40].

A GRACE update starts from a mini-batch $\mathcal{B} = \{(x_b, y_b)\}_{b=1}^B$ and a frozen teacher pool $\mathcal{T} = \{T_i\}_{i=1}^N$. The student first produces the logits, hidden states and supervised cross-entropy loss for the answer tokens. The router then uses the pooled student hidden state to select a sparse routed set $\mathcal{R}_b$ for each sample. For every selected teacher, trie-Wasserstein KD computes a per-teacher loss $\mathcal{L}_{KD}^{(b,i)}$, using the shared byte-trie geometry to compare student and teacher output distributions across different tokenizers. Gradient-agreement refinement then filters and reweights the routed teachers by comparing each teacher-specific KD gradient with the supervised CE gradient. At a high level, the routed KD signal for sample b is the weighted combination of teacher-specific losses $\tilde{\mathcal{L}}_{KD}^{(b)} = \sum_{i \in \mathcal{R}_b} w_i^{(b)} \mathcal{L}_{KD}^{(b,i)}$, where the final weights $w_i^{(b)}$ are determined by both the router and the gradient-agreement refinement step. This decomposition separates teacher selection from teacher validation: routing proposes teachers that appear relevant to the sample, while gradient agreement checks whether their current KD directions are compatible with the supervised learning objective. The resulting KD term is then combined with CE and router regularization in one optimization objective. The main notation used throughout the method is listed in Table 3.1.

Table 3.1: Main notation.

Symbol	Meaning
$\mathcal{D} = \{(x_n, y_n)\}_{n=1}^{N_{\text{data}}}$	Multimodal samples and target answer sequences.
$S_\theta, \mathcal{T} = \{T_i\}_{i=1}^N$	Trainable student VLM and frozen teacher pool.
G_ϕ	Trainable teacher routing network with parameters ϕ .
T_{train}	Total number of training steps.
$\mathbf{z}_s^{(b)}, \mathbf{z}_t^{(b,i)}$	Student and teacher logits for sample b .
$\mathbf{H}_s, \mathbf{h}$	Student hidden states and pooled routing feature.
$\rho_{\text{trie}}, \eta_{\text{boundary}}, K_{\text{TW}}, e_{\text{tail}}$	Trie edge decay, boundary-edge weight, sparse support size and residual tail edge.
k_{route}	Configured router top- k .
$\mathbf{r}, \mathbf{p}$	Router logits and soft router weights.
$\mathcal{C}_b, \mathcal{R}_b, \hat{\mathbf{p}}^{(b)}$	Top- k candidate teacher set, routed teacher set and renormalized routed weights.
$s_i, \hat{s}_i, \tilde{a}_i, a_i^{(b)}, \mathcal{A}_b$	Raw agreement score, EMA-smoothed agreement score, teacher active flag, sample active flag and active routed set.
$\omega^{(b)}, \mathbf{w}^{(b)}$	Score-softmax weights and final KD weights.
$\Omega_b^{CE}, \Omega_{b,i}^{KD}$	Supervised CE positions and matched KD positions.
$\mathbf{g}_{CE}, \mathbf{g}_i$	CE and teacher-KD parameter gradients used by GRACE.
$\beta_{\text{EMA}}, \beta_{\text{score}}, \epsilon_{\text{num}}, \epsilon_{\text{fb}}, \lambda_{\text{blend}}$	EMA decay, score-softmax sharpness, cosine stabilizer, fallback threshold and router-score blend exponent.
TopKRoute	Sparse top- k teacher routing with routed-weight renormalization.
Uniform fallback	Uniform weights over the active routed set.
Blend	Mixing routed weights with GRACE score weights.
EMA	Exponential moving average for GRACE scores.

3.2 Trie-Wasserstein Cross-Tokenizer Distillation

3.2.1 Supervised Answer Tokens

Distillation is applied only on **supervised answer tokens**, not on prompt tokens. For sample b , let Ω_b^{CE} denote the student positions whose labels are not ignored. For teacher T_i , let $\Omega_{b,i}^{KD}$ denote the matched answer positions after masking ignored labels on both student and teacher sides and truncating both spans to a common length. This is a position-wise approximation across tokenizers: after answer-span preprocessing, the j -th student and teacher answer positions are treated as corresponding positions. Trie-Wasserstein matching handles the token-distribution mismatch at each aligned position, while sequence-level byte alignment remains outside the objective. Restricting distillation to answer positions avoids wasting capacity on prompt formatting, padding and instruction wording, where teacher and student distributions carry no task-relevant signal.

3.2.2 Shared Byte Trie

Direct token-index matching assumes that the k -th coordinate of the teacher distribution and the k -th coordinate of the student distribution refer to the same decoded token. This assumption fails when tokenizers differ: the same vocabulary index may correspond to unrelated token strings in the teacher and student vocabularies. GRACE maps decoded teacher and student token strings into a **shared UTF-8 byte trie**, so tokens live in a shared prefix tree, not arbitrary vocabulary-index space. Tokens with common byte prefixes share the same initial trie path even if their tokenizer IDs are unrelated, allowing probability mass from different vocabularies to be compared through shared decoded-string structure. For each sample b , token position $j \in \Omega_{b,i}^{KD}$ and teacher T_i , let $\mathbf{z}_{s,j}^{(b)} \in \mathbb{R}^{V_s}$ be the student logits and $\mathbf{z}_{t,j}^{(b,i)} \in \mathbb{R}^{V_i}$ be the teacher logits. The student and teacher probability distributions are

$$\mathbf{p}_{s,j}^{(b)} = \text{softmax}\left(\frac{\mathbf{z}_{s,j}^{(b)}}{\tau_s}\right), \quad \mathbf{p}_{t,j}^{(b,i)} = \text{softmax}\left(\frac{\mathbf{z}_{t,j}^{(b,i)}}{\tau_t}\right), \quad (3.1)$$

where τ_s and τ_t are the student and teacher temperatures. In the sparse computation, ignored token IDs are masked before normalization so special and extended tokens do not receive probability mass in the trie.

Before insertion, each tokenizer piece is canonicalized into lexical content and optional boundary metadata. Visible boundary markers such as GPT-2-style space markers, SentencePiece space markers, WordPiece `##` markers and BPE `</w>` markers are stripped from the content string and recorded as boundary type. If the tokenizer decoder exposes leading whitespace directly, GRACE records that whitespace as a boundary signal and stores the stripped content bytes. The content is decoded, Unicode normalized and encoded as UTF-8 bytes. This keeps tokenizer conventions from splitting the shared content path at the root. Marker-only pieces remain explicit content pieces, so canonicalization never removes a token from the trie.

Each canonical content string is terminated with an EOS sentinel. The sentinel separates complete tokens from strict prefixes, so `doc` and `document` remain distinct leaves even though they share an initial path. If a tokenizer boundary marker was removed during canonicalization, GRACE adds a boundary edge after the EOS sentinel with weight η_{boundary} . This design shifts the comparison away from identical token IDs and keeps probability mass on similar decoded strings close, while treating whitespace and word-boundary conventions as small tokenizer-specific differences.

The trie defines a geometry over decoded token strings, illustrated in Figure 3.2 for the pairs `document/receipt` and `doc/rece`. That geometry is well suited to OCR-heavy VQA, where correct and near-correct answers often differ through prefixes, suffixes, spaces or number formats. For example, `<space>cat`, `cat` and `##cat` share the same content bytes for `cat`, with boundary edges recording tokenizer conventions separately. For a normal content-byte or EOS edge e , let $d(e)$ be its depth. GRACE assigns edge weights as follows:

$$w(e) = \begin{cases} \rho_{\text{trie}}^{d(e)-1}, & e \in E_{\text{content}} \cup E_{\text{EOS}}, \\ \eta_{\text{boundary}}, & e \in E_{\text{boundary}}. \end{cases} \quad (3.2)$$

where $0 < \rho_{\text{trie}} < 1$ and $\eta_{\text{boundary}} \geq 0$. Here E_{boundary} contains optional marker edges for space, continuation and end-of-word metadata. The first content edge from the root has weight 1 and deeper content edges decay geometrically. This makes shared prefixes cheaper to transport across than unrelated byte paths, while boundary edges keep space and continuation markers close to their underlying content. The loss therefore penalizes how far the student’s probability mass must move through the shared lexical tree to resemble the teacher’s mass, not whether the student assigns probability to the same token indices. The weighting still preserves a valid tree metric [17, 32] and different tokenizers may stop at different depths while still sharing the same byte-prefix edges whenever their decoded token strings overlap in byte-prefix structure.




Figure 3.2: Shared-trie tokenizer matching.

3.2.3 Trie-Wasserstein Objective

The shared byte trie induces a tree metric over decoded token strings. Let $\mathcal{V}_s$ and $\mathcal{V}_t$ denote the student and teacher token vocabularies being compared at a given answer position and let $\pi(v)$ denote the trie path of decoded token v . For an edge e , the student and teacher subtree masses are

$$m_s(e) = \sum_{\substack{v \in \mathcal{V}_s \\ e \in \pi(v)}} p_s(v), \quad m_t(e) = \sum_{\substack{u \in \mathcal{V}_t \\ e \in \pi(u)}} p_t(u).$$

$m_s(e)$ and $m_t(e)$ measure how much probability mass each distribution assigns to token paths passing through edge e . The W_1 optimal transport distance between teacher and student token distributions on the trie can then be written as a weighted sum of subtree-mass imbalances. The loss rewards lexical closeness, not exact token-ID equality:

$$W_T(p_s, p_t) = \sum_{e \in E} w(e) |m_s(e) - m_t(e)|. \quad (3.3)$$

This form follows from the standard W_1 Kantorovich optimal transport problem under the trie-induced tree metric [17, 32]. The full proof is provided in Appendix A. The sparse construction below applies the same formula to a bounded $\text{top-}K_{\text{TW}}$ support with a residual tail edge.

3.2.4 Sparse Support and Residual Tail Edge

Equation (3.3) is exact on the complete shared trie, but evaluating every vocabulary token at every answer position is expensive. GRACE builds a **sparse augmented trie**, so most useful distributional information, concentrated in high-probability tokens, is captured exactly. The tail edge handles the remainder: without it, omitted mass would silently disappear, making the teacher appear more confident than it really is and distorting the transport distance. At each aligned answer position, it masks ignored vocabulary entries, retains the **top- K_{TW} student and teacher tokens** and inserts those tokens through their canonical content paths, followed by EOS and any optional boundary edge. Probability mass outside the retained sets is not discarded. It is collapsed into a single **residual tail edge**. This gives an exact distance on the collapsed sparse trie and it approximates the complete trie distance by representing all omitted vocabulary mass with one tail atom. Let K_{TW} denote the sparse trie support size. Let $\mathcal{V}_s^*$ and $\mathcal{V}_t^*$ denote the non-ignored student and teacher vocabularies. Before constructing the sparse trie, the student and teacher probabilities are renormalized over the non-ignored vocabularies so that the retained and residual masses form valid distributions. Operationally, ignored entries are excluded from the softmax normalizer before top- K_{TW} selection. Let $\mathcal{K}_{s,j}^{(b)} \subset \mathcal{V}_s^*$ and $\mathcal{K}_{t,j}^{(b,i)} \subset \mathcal{V}_t^*$ be the retained token sets. The residual mass assigned to the tail edge is

$$r_{s,j}^{(b)} = 1 - \sum_{v \in \mathcal{K}_{s,j}^{(b)}} p_{s,j}^{(b)}(v), \quad r_{t,j}^{(b,i)} = 1 - \sum_{v \in \mathcal{K}_{t,j}^{(b,i)}} p_{t,j}^{(b,i)}(v). \quad (3.4)$$

Let $\pi(v)$ be the trie path for token v , including canonical content edges, the EOS edge and any optional boundary edge. Let e_{tail} be the global residual tail edge with $w(e_{\text{tail}}) = 1$. For any active trie edge e , the sparse subtree masses are

$$m_{s,j}^{(b)}(e) = \sum_{\substack{v \in \mathcal{K}_{s,j}^{(b)} \\ e \in \pi(v)}} p_{s,j}^{(b)}(v) + \mathbf{1}[e = e_{\text{tail}}] r_{s,j}^{(b)}, \quad (3.5)$$

$$m_{t,j}^{(b,i)}(e) = \sum_{\substack{v \in \mathcal{K}_{t,j}^{(b,i)} \\ e \in \pi(v)}} p_{t,j}^{(b,i)}(v) + \mathbf{1}[e = e_{\text{tail}}] r_{t,j}^{(b,i)}. \quad (3.6)$$

The token-level trie-Wasserstein loss is

$$\ell_{\text{TW}}(\mathbf{z}_{s,j}^{(b)}, \mathbf{z}_{t,j}^{(b,i)}) = \sum_{e \in \mathcal{E}_{\text{act},j}^{(b,i)}} w(e) \left| m_{s,j}^{(b)}(e) - m_{t,j}^{(b,i)}(e) \right|, \quad (3.7)$$

where $\mathcal{E}_{\text{act},j}^{(b,i)}$ contains the active canonical content edges, EOS edges, optional boundary edges and the residual tail edge from the student and teacher supports. Before the absolute value is applied, $m_{s,j}^{(b)}(e) - m_{t,j}^{(b,i)}(e)$ is positive when edge e carries more student mass and negative when it carries more teacher mass.




Figure 3.3: Sparse trie-Wasserstein support.

The per-sample, per-teacher KD loss is the mean over supervised answer tokens:

$$\mathcal{L}_{KD}^{(b,i)} = \frac{1}{|\Omega_{b,i}^{KD}|} \sum_{j \in \Omega_{b,i}^{KD}} \ell_{TW}(\mathbf{z}_{s,j}^{(b)}, \mathbf{z}_{t,j}^{(b,i)}). \quad (3.8)$$

Complexity analysis. The shared byte trie is built once with time and memory cost $O(V_s \bar{L}_s + V_t \bar{L}_t)$, where V_s and V_t are the student and teacher vocabulary sizes and $\bar{L}_s$ and $\bar{L}_t$ are the average canonical path lengths, including EOS and optional boundary edges. For one aligned answer position, extracting the top- K_{TW} student and teacher tokens costs $O(V_s \log K_{TW} + V_t \log K_{TW})$ with heap-based selection. Given the retained token sets, sparse subtree-mass accumulation walks only the retained canonical token paths and the residual tail edge, costing $O(K_{TW} \bar{L}_s + K_{TW} \bar{L}_t + A)$, where A is the number of active trie edges. If the implementation explicitly sorts or deduplicates active edge indices, this adds an $O(A \log A)$ term. Otherwise, indexed accumulation gives an $O(A)$ edge-aggregation cost. Runtime memory is $O(E + P_s + P_t + A)$, where E is the number of trie edges and P_s and P_t are the stored student and teacher path-index arrays. Unlike dense optimal transport, the sparse computation does not materialize a $|V_s| \times |V_t|$ transport matrix.




Figure 3.4: Teacher routing network.

3.3 Teacher Routing Network

3.3.1 Router Architecture

With multiple teachers, uniform averaging ignores sample-specific teacher strengths: a document question, a chart question and a scene-text question may each benefit from a different expert and the student’s own hidden states are the most direct signal of which kind of example is being processed. This makes GRACE routing different from sparse mixture-of-experts routing in standard transformers: MoE routing chooses which expert computes the forward pass, while GRACE routing chooses which frozen external teacher shapes the student’s learning signal. GRACE therefore attaches a **teacher routing network** to the linear language-model head, so the gate operates on the text hidden state immediately before token prediction. The full pipeline from pooled student features to sparse top- k teacher selection is shown in Figure 3.4.

Let $\mathbf{H}_s^{(b)}$ denote the student hidden states captured immediately before the language-model head for sample b . These are pooled over supervised answer positions to form the **routing feature**:

$$\mathbf{h}^{(b)} = \text{Pool}(\mathbf{H}_s^{(b)}) . \quad (3.9)$$

Pooling over answer positions captures what the student currently represents about the output, which is more informative than pooling over the full sequence. The routing feature is also student-state-dependent, not only dataset-dependent: the same image-question pair can produce a different routing decision as the student’s representations evolve during training, allowing teacher selection to adapt continuously. The operation uses the mean over supervised answer positions and requires at least one supervised answer token per sample. A learned SwiGLU-style MLP router $G_\phi(\cdot)$ maps the routing feature to gate logits,

$$\mathbf{r}^{(b)} = G_\phi(\mathbf{h}^{(b)}) . \quad (3.10)$$

These logits are converted to soft routing weights by softmax:

$$\mathbf{p}^{(b)} = \text{softmax}(\mathbf{r}^{(b)}) , \quad (3.11)$$




Figure 3.5: Top- k routing dynamics.

where $p_i^{(b)}$ is the initial gate weight for teacher T_i . The router then masks to the top- k_{route} teachers and renormalizes the selected weights. The internal stack uses LayerNorm, an up-projection, a gate-value split, SiLU gating and a down-projection, keeping score computation dense while teacher use remains sparse after top- k masking.

3.3.2 Top- k Routing

Soft routing weights are dense, but distillation should not use every teacher for every sample. GRACE converts them into a **sparse routed set** by selecting the highest-scoring teachers for sample b :

$$\mathcal{C}_b = \text{TopK}(\mathbf{r}^{(b)}, \min(k_{\text{route}}, N)). \quad (3.12)$$

The candidate set $\mathcal{C}_b$ serves as the routed teacher set $\mathcal{R}_b$. With $m_i^{(b)} = \mathbf{1}[i \in \mathcal{C}_b]$ denoting the top- k assignment mask, non-selected weights are zeroed and the remaining softmax weights are renormalized:

$$\hat{p}_i^{(b)} = \frac{m_i^{(b)} p_i^{(b)}}{\sum_{j=1}^N m_j^{(b)} p_j^{(b)}}. \quad (3.13)$$

Top- k routing makes supervision sparse and sample-specialized. Routing alone is not sufficient, however. A teacher can look relevant because its domain matches the input while its gradient still conflicts with the supervised target. For example, a document-specialized teacher may assign high probability to a plausible but incorrect field value. Gradient-agreement refinement, introduced in Section 3.4, addresses this residual failure. Figure 3.5 contrasts dense gate probabilities with the sparse routed assignment. The logged assignment rate records how often each teacher appears in the routed set.

3.3.3 Router Entropy and Z-Loss

Two auxiliary terms regularize teacher selection. **Entropy regularization** is implemented as a negative-entropy term that discourages premature routing collapse:

$$\mathcal{L}_{ent} = -\frac{1}{B} \sum_{b=1}^B H(\mathbf{p}^{(b)}). \quad (3.14)$$

The **router z-loss** keeps router logits numerically well behaved:

$$\mathcal{L}_z = \frac{1}{B} \sum_{b=1}^B \left(\log \sum_{i=1}^N e^{r_i^{(b)}} \right)^2. \quad (3.15)$$

Both terms are pure stabilizers: entropy discourages the router from collapsing onto a single teacher too early, while z-loss keeps routing logits numerically controlled. Neither competes with the CE or KD objectives.

3.4 Gradient-Agreement Refinement

Routed teachers may still conflict with the supervised objective after teacher selection. **Gradient-agreement refinement** addresses this by scoring each teacher’s KD direction against the supervised CE direction and using those scores inside the routed candidate set. Routing answers which teacher seems relevant to this sample. Gradient agreement answers which of those teachers is currently compatible with learning the correct answer. The two checks are complementary because domain relevance does not guarantee gradient compatibility, as established in Section 2.6.2.

3.4.1 Gradient Agreement Scoring

Let $\Theta \subseteq \theta$ denote the trainable student parameters used by the optimizer. To score each teacher, GRACE compares its **teacher-specific KD gradient** against the **supervised CE gradient** in student-parameter space. The CE gradient is

$$\mathbf{g}_{CE} = \nabla_{\Theta} \mathcal{L}_{CE}. \quad (3.16)$$

For teacher T_i , the KD gradient comes from that teacher’s mean distillation loss over the full mini-batch:

$$\mathbf{g}_i = \nabla_{\Theta} \left(\frac{1}{B} \sum_{b=1}^B \mathcal{L}_{KD}^{(b,i)} \right). \quad (3.17)$$

Here, $\mathcal{L}_{KD}^{(b,i)}$ is the trie-Wasserstein loss from Equation 3.8. These gradients score teacher compatibility while the final parameter update is still computed from the full training objective. As established in Section 2.6.2, a teacher whose KD gradient opposes the CE gradient pulls the student away from the correct answer even when the router selected it confidently.

3.4.2 Agreement Scores and Active Set

We quantify how well supervised learning and KD are directionally aligned using cosine similarity:

$$s_i = \frac{\langle \mathbf{g}_{CE}, \mathbf{g}_i \rangle}{\max(\|\mathbf{g}_{CE}\|_2 \|\mathbf{g}_i\|_2, \epsilon_{\text{num}})}. \quad (3.18)$$

The agreement score s_i is **teacher-level** within the current mini-batch and is broadcast across samples. Sample dependence enters through the routed set $\mathcal{R}_b$, routed weights $\hat{\mathbf{p}}^{(b)}$ and active set $\mathcal{A}_b$. To combine teacher-level agreement with sample-specific routing, GRACE first converts the agreement score into a binary teacher flag:

$$\tilde{a}_i = \mathbf{1}[s_i > \delta], \quad (3.19)$$

where δ is the GRACE threshold. This flag is then intersected with the routed teacher set for each sample:

$$a_i^{(b)} = \mathbf{1}[i \in \mathcal{R}_b] \tilde{a}_i, \quad \mathcal{A}_b = \{i \in \mathcal{R}_b : a_i^{(b)} = 1\}. \quad (3.20)$$

The intersection with $\mathcal{R}_b$ keeps the search inside the routed candidates: routing narrows the pool first and gradient agreement filters within it. The threshold δ is applied to the raw score s_i , while the EMA-smoothed score $\hat{s}_i$ controls score-softmax weighting and fallback spread. Agreement scores act as non-differentiable routing statistics, so the final optimization step does not backpropagate through the gradient-agreement computation itself or treat the scoring procedure as an additional trainable path.

3.4.3 Warmup, EMA Smoothing and Score Weights

Teacher weights can change abruptly when agreement scores are estimated from a single mini-batch. GRACE stabilizes this step with **warmup**, **EMA smoothing** and **score-softmax weighting**. Let $\rho_g \in [0, 1]$ denote the warmup ratio relative to the total number of training steps. During the first $\rho_g T_{\text{train}}$ steps, routed gate weights are left unchanged:

$$w_i^{(b)} = \hat{p}_i^{(b)}. \quad (3.21)$$

Early in training, the student’s representations and the router’s assignments are both still changing, so gradient-based filtering applied too soon risks rejecting useful teachers because the agreement scores are not yet reliable. Warmup prevents the agreement mechanism from overfitting to early gradient noise and missing genuine teacher usefulness. After warmup, routed teachers are refined with a smoothed score distribution. Let $\bar{s}_i^{\text{prev}}$ be the running exponential moving average (EMA) of past batch means for teacher T_i and let β_{EMA} be its decay. The smoothed score is

$$\hat{s}_i = \beta_{\text{EMA}} \bar{s}_i^{\text{prev}} + (1 - \beta_{\text{EMA}}) s_i. \quad (3.22)$$

EMA smoothing turns gradient agreement from a momentary vote into a reputation score. For nonempty active sets, smoothed scores are converted into a score-softmax weight vector over $\mathcal{A}_b$:

$$\omega_i^{(b)} = \frac{\exp(\beta_{\text{score}} \hat{s}_i) \mathbf{1}[i \in \mathcal{A}_b]}{\sum_{j=1}^N \exp(\beta_{\text{score}} \hat{s}_j) \mathbf{1}[j \in \mathcal{A}_b]}, \quad (3.23)$$

β_{score} controls how sharply the highest-scoring active teacher is preferred, setting up the fallback and blending cases below.

3.4.4 Weight Blending and Fallback

Final KD weights are assigned through three mutually exclusive cases for each sample. These cases prevent the refinement step from overreacting when agreement evidence is missing or weak.

No active teacher. If $\mathcal{A}_b = \emptyset$, the routed weights are retained:

$$w_i^{(b)} = \hat{p}_i^{(b)}. \quad (3.24)$$

Low-confidence scores. If active teachers exist but their smoothed scores are too close, uniform weights are used over the active set. The score range is

$$\Delta^{(b)} = \max_{i \in \mathcal{A}_b} \hat{s}_i - \min_{i \in \mathcal{A}_b} \hat{s}_i. \quad (3.25)$$

If $\mathcal{A}_b \neq \emptyset$ and $\Delta^{(b)} < \epsilon_{\text{fb}}$, then

$$w_i^{(b)} = \frac{\mathbf{1}[i \in \mathcal{A}_b]}{|\mathcal{A}_b|}. \quad (3.26)$$

Confident scores. Otherwise, router weights are blended with score weights. Let $\lambda_{\text{blend}} \in [0, 1]$ be the router-blend exponent. The unnormalized blended weights are

$$\tilde{w}_i^{(b)} \propto \left(\hat{p}_i^{(b)}\right)^{\lambda_{\text{blend}}} \left(\omega_i^{(b)}\right)^{1-\lambda_{\text{blend}}}, \quad (3.27)$$

followed by normalization over the active routed teachers:

$$w_i^{(b)} = \frac{\tilde{w}_i^{(b)} \mathbf{1}[i \in \mathcal{A}_b]}{\sum_{j=1}^N \tilde{w}_j^{(b)} \mathbf{1}[j \in \mathcal{A}_b]}. \quad (3.28)$$

Equations 3.23, 3.26 and 3.28 define the gradient-refinement operator: the router chooses $\mathcal{R}_b$, agreement filtering narrows it to $\mathcal{A}_b$ and final weights are retained from the router, assigned uniformly or geometrically blended depending on confidence. These regimes are shown in Figure 3.6.

This design keeps teacher selection sparse throughout: gradient-agreement refinement can narrow the routed set but never expands it, so KD cost remains bounded by k_{route} teachers per sample.



Figure 3.6: Representative KD weighting regimes.

3.5 Routed KD Aggregation and Final Objective

The final teacher weights $w_i^{(b)}$ combine the routed teacher set $\mathcal{R}_b$, router weights $\hat{\mathbf{p}}^{(b)}$ and gradient-agreement scores. The objective aggregates routed per-teacher KD losses into a sample-level term, averages those terms across the mini-batch and then adds supervised and router-regularization losses.

Sample-level aggregation. For sample b , only teachers in the routed set $\mathcal{R}_b$ contribute to the KD loss:

$$\tilde{\mathcal{L}}_{KD}^{(b)} = \sum_{i \in \mathcal{R}_b} w_i^{(b)} \mathcal{L}_{KD}^{(b,i)}. \quad (3.29)$$

Non-routed teachers have zero contribution, while routed teachers are weighted by the fallback, uniform-active or blended rule. All post-routing refinement decisions enter the optimization through $w_i^{(b)}$.

Batch-level aggregation. The mini-batch KD term averages the sample-level KD losses:

$$\mathcal{L}_{KD} = \frac{1}{B} \sum_{b=1}^B \tilde{\mathcal{L}}_{KD}^{(b)}. \quad (3.30)$$

This scalar term supplies teacher supervision to the final objective. It stays compatible with heterogeneous teachers because each per-teacher loss uses trie-Wasserstein matching, which is tokenizer-agnostic and avoids direct token-index KL divergence.

Final objective. The complete GRACE objective is

$$\mathcal{L} = \mathcal{L}_{CE} + \alpha_{KD} \mathcal{L}_{KD} + \lambda_{ent} \mathcal{L}_{ent} + \lambda_z \mathcal{L}_z, \quad (3.31)$$

The terms in this objective follow the dependency chain: $\mathcal{L}_{CE}$ is the compass that fixes the learning direction, $\mathcal{L}_{KD}$ provides the teacher map, routing chooses which map to consult and gradient agreement checks whether the map currently points in the compass direction. The coefficient α_{KD} controls the strength of teacher supervision, while λ_{ent} and λ_z control the auxiliary router losses. Only the student parameters θ and router parameters ϕ are optimized, while all teacher parameters remain frozen.

Chapter 4

Evaluation Methodology

This chapter describes the experimental setup used to evaluate GRACE and all compared methods. It covers the student models, benchmarks, training configuration, evaluation metrics and baselines. The ablation settings and implementation details needed to reproduce the main results are included at the end of the chapter.

4.1 Experimental Setup

All experiments follow a shared evaluation protocol for fair comparison across methods. The experiments instantiate the student side with **SmolVLM-500M** and **SmolVLM-256M**, evaluated on multimodal benchmarks [15]. All models are evaluated with **VLM-Eval-Kit** [44], which provides a shared inference, post-processing and metric-computation pipeline for the three benchmarks. VLM-Eval-Kit uses a generation-based evaluation interface. For each benchmark sample, the toolkit loads the image-question pair, applies the dataset-specific prompt construction logic and calls the model’s generation or chat interface to produce a free-form text answer. This choice is important for VLMs because it evaluates the same interface used at inference time and avoids replacing answer generation with likelihood scoring over a fixed set of options. After generation, the toolkit writes model predictions to result files and invokes the benchmark-specific evaluator.

The post-processing stage depends on the task format. For structured yes/no or multiple-choice tasks, VLM-Eval-Kit can use exact matching or an LLM-based choice extractor to recover the final option from the generated text. For open-answer VQA benchmarks such as the ones used here, the generated answer is normalized before scoring so that superficial formatting differences do not dominate the metric. The final score is then computed with the official metric associated with each dataset. For example, a TextVQA prediction is matched against the human answer set using VQA accuracy, so a short answer such as **stop** is credited according to annotator agreement. A DocVQA prediction is scored with ANLS, so a near miss such as **INV-4572** versus **4572** receives partial credit only when the normalized edit similarity is high enough. A ChartQA numeric answer is evaluated with relaxed accuracy, so a value within 5% of the reference number is counted as correct. In this thesis, using the same toolkit for all methods fixes the prompt template, decoding interface, answer normalization and metric implementation, so evaluation-code differences do not drive the reported performance gaps.



Figure 4.1: Representative dataset examples.¹

4.2 Datasets

We evaluate three text-centric VQA benchmarks: **TextVQA**, **DocVQA** and **ChartQA**. These benchmarks test whether the model can read visual text, understand questions and perform basic reasoning across natural images, documents and charts. They were chosen because they cover complementary visual domains and combine OCR-heavy understanding with layout interpretation and numerical reasoning [4–6].

TextVQA tests question answering on natural images containing scene text, requiring the model to read and reason over text embedded in real-world scenes [4]. *Sample question*: “What word is written on the stop sign?”

DocVQA tests question answering on scanned document images, where layout, paragraphs, tables and forms are central to answering correctly [5]. *Sample question*: “What is the invoice number shown in this document?”

ChartQA tests question answering over charts, combining chart reading with numerical and logical reasoning [6]. *Sample question*: “Which year has the highest sales value in the chart?”

Figure 4.1 shows the visual differences behind this benchmark choice. The TextVQA panel represents natural-image reading, where the relevant cue is embedded in an object in the scene. The DocVQA panel is a receipt with fields, item rows and totals arranged by document layout. The ChartQA panel is a bar chart, where answering requires comparing bar heights and reading values from the plot structure. The evaluation therefore covers natural scene text, document layouts and chart reasoning. The split sizes used throughout the evaluation chapter are listed in Table 4.1.

¹Images adapted from Wikimedia Commons sources [45–47].

Table 4.1: Dataset split sizes.

Dataset	Train	Val	Test	Total
TextVQA	34,602	5,000	5,734	45,336
DocVQA	39,463	5,349	5,188	50,000
ChartQA-H (human)	7,398	960	1,250	9,608
ChartQA-M (machine)	20,901	960	1,250	23,111
ChartQA (H+M, combined)	28,299	1,920	2,500	32,719

4.3 Training Configuration

Before training, all benchmark samples are converted into a unified **LLaVA-style multimodal instruction format** [3]. Each example is stored as an image plus a short user-assistant dialogue, as shown in Listing 4.1, so TextVQA, DocVQA and ChartQA can share the same data loader, prompt template and answer supervision interface. This format is useful because the model sees every task as the same conditional generation problem: given an image and a natural-language user query, generate the assistant answer. It also separates the visual input path from the textual instruction path while keeping the answer span explicit, which makes it easier to mask prompt tokens, compute supervised loss only on assistant tokens and align teacher logits with the same answer positions during distillation.

```

1 {
2   "image": "sample_000123.jpg",
3   "conversations": [
4     {"from": "human", "value": "<image>\nWhat is the invoice number?"},
5     {"from": "gpt", "value": "4572"}
6   ]
7 }
```

Listing 4.1: Representative LLaVA-style sample used after dataset conversion.

We run all training jobs on a single high-memory **NVIDIA H200** graphics processing unit (GPU) or on comparable single-GPU **Vast.ai** instances, depending on availability. Distillation training uses **DeepSpeed ZeRO Stage 2**, **Brain Floating Point 16 (bf16)** mixed precision, **TensorFloat-32 (TF32)** matrix multiplication and gradient checkpointing to control memory use while keeping the student fully trainable and all teachers frozen [48, 49]. Full optimization and routing parameter values are reported in Appendix B, Table B.1. Teacher-logit storage and remote caching are documented in Appendix B.2 and Appendix B.3 for reproducibility purposes.

4.4 Evaluation Metrics

We report benchmark-specific official metrics for each dataset.

TextVQA. For TextVQA, we report the official VQA Accuracy metric. For a predicted answer $\hat{a}$, the per-question score is

$$s_{\text{VQA}} = \min \left(\frac{\# \text{annotators that provided } \hat{a}}{3}, 1 \right),$$

and the final score is the mean over all questions [4].

DocVQA. We report ANLS, the Average Normalized Levenshtein Similarity. The final ANLS is the average score over all samples [5]:

$$\text{NLS}(a, \hat{a}) = 1 - \frac{d_{\text{Lev}}(a, \hat{a})}{\max(|a|, |\hat{a}|)}, \quad s_{\text{ANLS}} = \begin{cases} \text{NLS}(a, \hat{a}), & \text{NLS} \geq 0.5, \\ 0, & \text{otherwise.} \end{cases}$$

ChartQA. For ChartQA, we report Relaxed Accuracy [6]. A prediction is correct if it matches exactly for text answers. For numeric answers, the score handles the zero-target case explicitly:

$$s_{\text{ChartQA}} = \begin{cases} \mathbf{1}[\hat{y} = y], & y = 0, \\ \mathbf{1}\left[\frac{|\hat{y} - y|}{|y|} \leq 0.05\right], & y \neq 0. \end{cases}$$

4.5 Baselines and Compared Methods

For consistency with Chapter 5, we organize the compared methods into four groups: non-distillation references, same-family distillation baselines, different-family single-teacher baselines and heterogeneous teacher baselines. Before introducing the method-specific baselines, the role of each baseline family is grouped in Table 4.2. The grouping separates student-only references, same-family KD, tokenizer-mismatched single-teacher transfer and heterogeneous teacher methods.

The non-distillation references consist of the original student and the full-finetuning baseline, which provide student-only reference points. The same-family baselines use **SmolVLM2-2.2B-Instruct** as teacher for forward KL, reverse KL and Jensen-Shannon distillation, isolating standard teacher-student transfer when the tokenizer and output space are matched [11, 15, 29, 50, 51]. The different-family single-teacher setting uses ULD with a teacher outside the SmolVLM family, testing tokenizer-mismatched transfer without teacher routing [16]. The heterogeneous teacher setting compares uniform averaging, SHM-CKA, REINFORCE and GRACE over a mixed teacher pool, separating static averaging, hidden-state matching and adaptive teacher assignment [17, 21, 22, 31].

Table 4.2: Baseline groups.

Group	Methods	Teacher setup	Purpose
Non-distillation references	Original, full finetuning	No teacher	Student-only reference points
Same-family baselines	Forward KL, Reverse KL, Jensen-Shannon	SmolVLM2-2.2B-Instruct	Standard single-teacher objectives within the SmolVLM family
Cross-family single-teacher	ULD	Qwen2-VL-2B-Instruct	Cross-family transfer without routed teacher selection
Heterogeneous teacher baselines	Uniform averaging, SHM-CKA, REINFORCE, GRACE (Ours)	Mixed teacher set from multiple VLM families	Static averaging vs. hidden-state matching vs. adaptive teacher assignment

For the trainable competitor methods, we keep the supervised term $\mathcal{L}_{CE}$. The logits-only baselines differ only in the KD signal, but SHM-CKA also adds a hidden-state alignment term and REINFORCE also adds a REINFORCE-style policy term [19, 20, 52]. In the descriptions below, γ denotes a generic KD mixing coefficient. α_{KD} is reserved for the GRACE training objective as defined in Equation 3.31.

Original. Directly evaluates the released student without additional optimization, thereby measuring the native capability of the student model before task adaptation or any teacher-guided distillation signal.

Full finetuning. Optimizes only $\mathcal{L}_{CE}$ and uses no teacher, which isolates the effect of supervised task adaptation from any teacher-guided distillation signal.

Forward KL. Uses $\mathcal{L}_{KD} = \text{KL}(p_t \| p_s)$ for same-family soft-target distillation when teacher and student vocabularies are aligned [11].

Reverse KL. Uses $\mathcal{L}_{KD} = \text{KL}(p_s \| p_t)$ [51, 53] as a student-centered divergence baseline, which tests whether matching the teacher from the student’s current distribution is suitable for VLM distillation. Because the reverse-KL gradient is weighted by p_s , this objective tends to emphasize regions where the student already assigns probability.

Jensen-Shannon. Uses the symmetric divergence

$$\frac{1}{2}\text{KL}(p_s \| m) + \frac{1}{2}\text{KL}(p_t \| m), \quad m = \frac{1}{2}(p_s + p_t).$$

It tests whether a bounded symmetric objective is more stable than KL [29, 51].

ULD. Uses $\mathcal{L}_{KD} = \|\pi(p_s) - \pi(p_t)\|_1$ [16]. This is the main tokenizer-mismatched logits baseline because it compares ranked probability mass, not token IDs.

SHM-CKA. Extends ULD with a hidden-state CKA penalty [16, 19, 31]. For the heterogeneous setting, both the ULD and layer-matching terms are averaged uniformly across teachers before combining with $\mathcal{L}_{CE}$. Full implementation details are given in Appendix B.1.

Uniform multi-teacher. Assigns fixed equal weights to all teachers in the heterogeneous pool, $\mathcal{L}_{KD} = \frac{1}{M} \sum_{i=1}^M \mathcal{L}_{KD}^{(i)}$ [20, 21], thereby testing whether teacher diversity alone is sufficient without adaptive routing or gradient-aware weighting.

REINFORCE. Replaces deterministic routing with a stochastic policy that samples teachers and updates the selector with REINFORCE [20, 52], providing an adaptive routing baseline that does not use the gradient-agreement refinement introduced by GRACE. The selector is trained with the policy-gradient term

$$\mathcal{L}_{\text{pol}} = -(r - b) \log p(\mathbf{m} \mid \boldsymbol{\pi}), \quad (4.1)$$

where r is the reward, b is an EMA reward baseline and $\mathbf{m}$ is the sampled teacher mask. Its final training objective is

$$\mathcal{L}_{\text{RS}} = \mathcal{L}_{CE} + \gamma \mathcal{L}_{KD} + \lambda_{\pi} \mathcal{L}_{\text{pol}}. \quad (4.2)$$

Additional baseline implementation details are reported in Appendix B.1.

4.6 Ablation Settings

The ablation study in Section 5.3 isolates the main components of GRACE on DocVQA while keeping the teacher set and training protocol close to the final method. The variants change one mechanism at a time so that the result table can separate teacher diversity, sparse routing, gradient-agreement weighting and EMA score smoothing. The effect of EMA smoothing is isolated in the $\beta_{\text{EMA}} = 0$ ablation in Section 5.3.

Cross-family ULD. The single-teacher cross-family baseline uses one teacher outside the SmolVLM family and optimizes

$$\mathcal{L} = \mathcal{L}_{CE} + \alpha_{\text{KD}} \mathcal{L}_{ULD}, \quad \mathcal{L}_{ULD} = \|\pi(p_s) - \pi(p_t)\|_1.$$

Uniform averaging. The uniform variant adds the full heterogeneous teacher pool but assigns fixed equal weights:

$$\mathcal{L} = \mathcal{L}_{CE} + \alpha_{\text{KD}} \frac{1}{M} \sum_{i=1}^M \mathcal{L}_{TW}^{(i)}.$$

This setting tests whether teacher diversity alone is enough without adaptive selection.

Routing only. The routing-only variant keeps sparse teacher selection but replaces post-routing weights with uniform weights within the routed set:

$$w_i^{(b)} = u_i^{(b)} = \frac{\mathbf{1}[i \in \mathcal{R}_b]}{|\mathcal{R}_b|}.$$

This setting tests whether top- k routing is useful without gradient-agreement refinement.

Full GRACE. The full method combines trie-Wasserstein KD, sparse routing and gradient-agreement weighting:

$$w_i^{(b)} \propto (\hat{p}_i^{(b)})^{\lambda_{\text{blend}}} (\omega_i^{(b)})^{1-\lambda_{\text{blend}}}.$$

Two additional ablation variants isolate the final weighting mechanism, reported in Table 5.6. The $\lambda_{\text{blend}} = 1.0$ variant keeps the active-set threshold and fallback rules but removes the score-softmax factor from the confident-score case:

$$w_i^{(b)} \propto \hat{p}_i^{(b)} \mathbf{1}[i \in \mathcal{A}_b].$$

This row is an active-set-restricted router-weight ablation, not a pure router-only ablation over all routed teachers. The $\beta_{\text{EMA}} = 0$ variant removes temporal smoothing by setting $\hat{s}_i = s_i$ before score-softmax weighting.

The training runs reported in this thesis rely on an offline teacher-logit preparation step and a remote logit-streaming infrastructure that are separate from the fair-comparison protocol described above. Teacher logits are generated once and stored on Cloudflare R2. During training, they are streamed lazily through a 100 GB local cache without copying the full logit store locally. Full engineering details, including the caching script, path conventions, streaming loader and the logit-streaming settings table, are given in Appendix B.2 and Appendix B.3.

Chapter 5

Experiment Results

This chapter presents the empirical results for all experiments described in Chapter 4. The main comparison evaluates GRACE against non-distillation references, same-family logits baselines, cross-family single-teacher transfer and heterogeneous multi-teacher methods across TextVQA, DocVQA and ChartQA. The ablation study then isolates the contribution of each component and the efficiency analysis addresses training and storage costs under the reported setup.

5.1 Main Comparison Results

This section reports the main empirical comparison used to answer RQ1. The results are organized from simpler reference points to the full heterogeneous multi-teacher setting: original students, supervised fine-tuning, same-family logits KD, cross-family ULD and multi-teacher distillation methods. The section first establishes the student and teacher reference ranges, then compares all methods on TextVQA, DocVQA and ChartQA. This ordering separates two questions for the reader: whether distillation improves over the student-only reference point and whether adaptive multi-teacher weighting adds value beyond fixed averaging or single-teacher transfer.

5.1.1 Baseline Results

Before applying adaptation or distillation, we report the original performance of the two student models, **SmolVLM-500M** and **SmolVLM-256M**, using **VLM-Eval-Kit** [44] under the evaluation pipeline described in Chapter 4. The student-only starting point for all later fine-tuning and distillation comparisons is given in Table 5.1. The teacher-side reference range is reported in Table 5.2. The gap between these student models and the higher-scoring teachers motivates the distillation experiments that follow.

Table 5.1: Original SmolVLM student results.

Model	TextVQA	DocVQA	ChartQA
SmolVLM-500M	56.75	59.52	62.80
SmolVLM-256M	47.17	48.64	45.04

Table 5.2: Teacher and reference-model results.

Teacher Model	TextVQA	DocVQA	ChartQA
Qwen2.5-VL-3B-Instruct	79.30	93.90	94.24
InternVL3-1B	74.10	81.90	75.30
Gemma-3-4B-it	57.80	75.80	63.60
Qwen2-VL-2B-Instruct	79.70	90.10	89.28

Using the highest teacher score on each benchmark as the reference, the original teacher-student gap is 22.95, 34.38 and 31.44 points for SmolVLM-500M on TextVQA, DocVQA and ChartQA. For SmolVLM-256M, the corresponding gaps are 32.53, 45.26 and 49.20 points. These gaps define the headroom used later in the gap-to-teacher analysis.

5.1.2 Overall Comparison

The main benchmark results across original models, direct fine-tuning, standard logits distillation baselines and GRACE are collected in Table 5.3. This main result table reports the highest-scoring completed run for each method and benchmark. Uniform, SHM-CKA, REINFORCE and GRACE use a four-teacher mixed set spanning the Qwen, InternVL and Gemma families. Chapter 4 documents the four-teacher configuration used across all multi-teacher experiments. We report the baselines in four groups. The non-distillation references are the original student and full-finetuning runs. The same-family baselines use **SmolVLM2-2.2B-Instruct** as teacher for forward KL, reverse KL and Jensen-Shannon divergence [11, 12, 31]. The different-family single-teacher setting uses ULD with a teacher outside the SmolVLM family [16]. The multi-teacher baselines compare four-teacher Uniform averaging with adaptive methods such as SHM-CKA, REINFORCE and GRACE [20–22, 52].

Table 5.3: Main comparison across TextVQA, DocVQA and ChartQA.

Category	Method	Teacher	Student	TextVQA	DocVQA	ChartQA
Non-distillation references						
Original	None	None	SmolVLM-500M	56.75	59.52	62.80
			SmolVLM-256M	47.17	48.64	45.04
Reference	Finetuning	None	SmolVLM-500M	56.96	61.96	63.00
			SmolVLM-256M	49.83	42.03	49.68
Same-family distillation baselines						
Same-family	Forward KL	SmolVLM2-2.2B	SmolVLM-500M	59.00	63.55	64.52
			SmolVLM-256M	50.88	49.46	56.80
	Reverse KL		SmolVLM-500M	44.73	29.28	30.16
			SmolVLM-256M	29.75	23.28	17.20
	Jensen-Shannon		SmolVLM-500M	42.33	36.43	43.75
			SmolVLM-256M	33.84	33.12	50.93
Different-family single-teacher baselines						
Cross-family	ULD	Qwen2-VL-2B-Instruct	SmolVLM-500M	61.49	65.28	66.92
			SmolVLM-256M	53.95	52.24	60.96
Multi-teacher baselines						
Multi-teacher	Uniform	†	SmolVLM-500M	63.88	69.59	75.60
			SmolVLM-256M	54.57	55.70	63.68
	SHM-CKA		SmolVLM-500M	60.47	63.32	76.88
			SmolVLM-256M	55.93	57.24	62.40
	REINFORCE		SmolVLM-500M	63.26	64.00	70.27
			SmolVLM-256M	56.31	58.53	68.74
	GRACE (Ours)		SmolVLM-500M	69.38	73.11	80.20
			SmolVLM-256M	60.73	60.48	72.32

† Teacher pool: Qwen2-VL-2B-Instruct, Qwen2.5-VL-3B-Instruct, InternVL3-1B, Gemma-3-4B-it.

On DocVQA, GRACE rises from 69.59 to 73.11 on SmolVLM-500M and from 55.70 to 60.48 on SmolVLM-256M relative to the Uniform baseline. ChartQA shows the largest margins over Uniform: 75.60 to 80.20 for SmolVLM-500M and 63.68 to 72.32 for SmolVLM-256M. Several score patterns are consistent across the three benchmarks. Forward KL gives small gains over the original students, while reverse KL reduces performance and Jensen-Shannon varies by task. Cross-family ULD scores above the same-family logits losses on all three tasks, indicating that tokenizer-mismatched distillation is more effective in this setting than vocabulary-matched KL. Uniform averaging adds another increase over ULD, with gains of +2.39, +4.31 and +8.68 for SmolVLM-500M on TextVQA, DocVQA and ChartQA. The additional improvement from single-teacher ULD to GRACE comes from trie-Wasserstein KD, sample-dependent sparse routing and gradient-agreement refinement after routing. GRACE also improves on the uniform multi-teacher baseline for both student sizes. The gain is smaller on TextVQA and DocVQA than on ChartQA, indicating that adaptive weighting has the largest observed benefit when the task mixes visual reading, numerical reasoning and teacher specialization across datasets.




Figure 5.1: Grouped benchmark comparison.

Figure 5.1 summarizes the main ranking by showing that GRACE scores above the highest non-GRACE baseline across all three benchmarks and both student sizes.

To provide qualitative evidence of attention improvement, we use DEX-AR attribution maps [54]. For token t , layer l and head h , DEX-AR forms a gradient-based attribution map

$$\mathbf{G}_t^{(l,h)} = \nabla_{\mathbf{A}^{(l,h)}} y_t \odot \mathbf{A}^{(l,h)},$$

then keeps visually grounded heads with

$$w^{(l,h)} = \max(S_{\text{img}}^{(l,h)} - S_{\text{text}}^{(l,h)}, 0),$$

and aggregates

$$\mathbf{H}_t = \sum_{l,h} w^{(l,h)} \mathbf{G}_t^{(l,h)}, \quad \mathbf{H}_{\text{sent}} = \sum_t \delta_t \mathbf{H}_t.$$

Here $\delta_t = 1/|\Omega_{\text{ans}}|$ for answer tokens $t \in \Omega_{\text{ans}}$, where Ω_{ans} is the set of generated answer token positions and $\delta_t = 0$ otherwise. The sentence-level map is therefore the uniform average of token-level maps over the generated answer span. Raw attention maps and generic saliency methods are not well suited here because the model is autoregressive and the quantity of interest is which image regions drove the answer, not which regions are visually salient in general. DEX-AR is designed for this target: it attributes each answer token back to the input image and averages those attributions across the answer span. Figure 5.2 summarizes the pipeline before the qualitative TextVQA example, making the later heatmaps easier to interpret as answer-grounding evidence, not generic saliency maps.




Figure 5.2: DEX-AR overview.

Figure 5.3 gives a representative TextVQA example. Sample 3 asks “*who is this copyrighted by?*” The relevant visual evidence is the copyright text on the lower-left border. The original SmolVLM-500M assigns high attention to the *Virgin* logo, not this region. GRACE shifts attention toward the copyright text, reducing this error.



(a) Input image



(b) Original SmolVLM-500M



(c) GRACE SmolVLM-500M

Figure 5.3: Representative TextVQA attention example.

Answer to RQ1

GRACE improves over direct finetuning, same-family logits KD, ULD and the compared heterogeneous multi-teacher baselines in Table 5.3. The score order places same-family logits baselines below cross-family ULD, cross-family ULD below most multi-teacher settings and GRACE above every compared baseline. Within the multi-teacher group, GRACE scores above SHM-CKA, REINFORCE and Uniform for both student sizes across the three benchmarks. The largest gains come from combining the trie-Wasserstein loss with gradient-based alignment. Table 5.5 quantifies how much of the teacher-student gap is closed by GRACE.





Figure 5.4: Uniform and GRACE training curves.

Figure 5.4 shows that the GRACE run produces a smoother and more stable training curve than Uniform on DocVQA. The Uniform curve shows sharper oscillations, while the GRACE curve decreases more steadily after the early training phase.



(a) Without GRACE refinement



(b) With GRACE

Figure 5.5: Routed-teacher load examples.

Figure 5.5 contrasts collapsed and stable routing behavior. Without the full GRACE recipe, the routed load concentrates on one teacher, reducing teacher diversity. With GRACE, the load stays distributed across the teacher pool, so the student continues to receive heterogeneous supervision. The relative gains in Table 5.4 are consistent with this routing behavior. Relative gain is computed as $(\text{GRACE} - \text{baseline})/\text{baseline}$.

Table 5.4: Relative gains of GRACE.

Benchmark	vs Original	vs Finetuning	vs Uniform	vs Highest non-GRACE
<i>SmolVLM-500M</i>				
TextVQA	+12.63 (+22.26%)	+12.42 (+21.80%)	+5.50 (+8.61%)	+5.38 (+8.41%)
DocVQA	+13.59 (+22.83%)	+11.15 (+18.00%)	+3.52 (+5.06%)	+3.52 (+5.06%)
ChartQA	+17.40 (+27.71%)	+17.20 (+27.30%)	+4.60 (+6.08%)	+3.32 (+4.32%)
<i>SmolVLM-256M</i>				
TextVQA	+13.56 (+28.75%)	+10.90 (+21.88%)	+6.16 (+11.29%)	+4.42 (+7.85%)
DocVQA	+11.84 (+24.34%)	+18.45 (+43.90%)	+4.78 (+8.58%)	+1.95 (+3.33%)
ChartQA	+27.28 (+60.57%)	+22.64 (+45.57%)	+8.64 (+13.57%)	+3.58 (+5.21%)

The gains over the original SmolVLM-500M student are large: +22.26% on TextVQA, +22.83% on DocVQA and +27.71% on ChartQA. The gains over direct finetuning are almost as large, at +21.80%, +18.00% and +27.30%, respectively. The gain over static uniform averaging is smaller, but it remains consistent across all three tasks, isolating the value of adaptive teacher weighting from the value of simply using more than one teacher. The same trend holds for SmolVLM-256M, with relative gains of +28.75%, +24.34% and +60.57% on the three benchmarks. For the gap-to-teacher analysis, let the reference point be the highest-scoring teacher in the GRACE comparison pool reported in Table 5.2. Let $g_{\text{orig}} = y_T - y_{\text{orig}}$ and $g_{\text{grace}} = y_T - y_{\text{grace}}$. The fraction of the gap closed is $1 - g_{\text{grace}}/g_{\text{orig}}$.

Table 5.5: Gap-to-teacher analysis for SmolVLM-500M.

Benchmark	Highest teacher	Original gap	GRACE gap	Gap closed
TextVQA	79.70	22.95	10.32	55.03%
DocVQA	93.90	34.38	20.79	39.53%
ChartQA	94.24	31.44	14.04	55.34%

The gain is not merely an artifact of comparison with lower-scoring baselines. On two benchmarks, GRACE removes about half of the original teacher-student gap and even on DocVQA, where the teacher ceiling is highest, it still closes almost 40% of the gap. **The distilled student can outperform individual teachers.** On TextVQA, GRACE reaches 69.38 and exceeds Gemma-3-4B-it at 57.80 by +11.58 points (+20.03%). On ChartQA, it exceeds InternVL3-1B at 75.30 by +4.90 points (+6.51%) and Gemma-3-4B-it at 63.60 by +16.60 points (+26.10%). This shows that routed multi-teacher supervision can combine useful signals that are not available from any single teacher alone.

One negative result is also consistent. Reverse KL scores much lower than forward KL. This can be explained by the different student-logit gradients. Let $p_s = \text{softmax}(z_s)$ be the student distribution and let p_t be the fixed teacher distribution. For forward KL,

$$\mathcal{L}_{FKL} = \text{KL}(p_t \| p_s) = \sum_v p_t(v) \log \frac{p_t(v)}{p_s(v)}, \quad \frac{\partial \mathcal{L}_{FKL}}{\partial z_s(v)} = p_s(v) - p_t(v).$$

Thus, every teacher-supported token directly pulls the student distribution toward the teacher’s soft target [11]. Reverse KL gives

$$\mathcal{L}_{RKL} = \text{KL}(p_s \| p_t) = \sum_v p_s(v) \log \frac{p_s(v)}{p_t(v)}.$$

The student-logit gradient is

$$\frac{\partial \mathcal{L}_{RKL}}{\partial z_s(v)} = p_s(v) \left(\log \frac{p_s(v)}{p_t(v)} - \mathcal{L}_{RKL} \right).$$

The key difference is the factor $p_s(v)$. If $p_t(v) > 0$ but $p_s(v) \rightarrow 0$, then

$$\frac{\partial \mathcal{L}_{FKL}}{\partial z_s(v)} \rightarrow -p_t(v), \quad \frac{\partial \mathcal{L}_{RKL}}{\partial z_s(v)} \rightarrow 0.$$

Teacher-supported tokens that the student currently assigns low probability therefore receive lower-magnitude correction under reverse KL. In finite training, this student-weighted gradient can make reverse KL less useful for answer-token distillation, especially when teacher distributions contain useful secondary mass over OCR variants, numeric formats or plausible answer alternatives. We interpret this as an optimization-bias explanation, not as a universal rule, because recent LLM distillation work shows that the simple mode-seeking interpretation of reverse KL does not always fully describe KD behavior [55].

SHM-CKA alone is also not sufficient because representation similarity is not the same as task-useful output supervision. CKA measures hidden-state geometry [19], but hidden-state similarity does not imply output equivalence:

$$\text{CKA}(H_s, H_t) \approx 1 \not\Rightarrow p_s(\cdot | x) \approx p_t(\cdot | x).$$

A simple linear-head example shows why. Suppose the teacher representation is an orthogonal rotation of the student representation, $H_t = H_s R$ with $R^\top R = I$. Linear CKA gives

$$\text{CKA}(H_s, H_t) = \frac{\|H_t^\top H_s\|_F^2}{\|H_t^\top H_t\|_F \|H_s^\top H_s\|_F} = 1.$$

However, the answer logits can still differ:

$$z_s = H_s W_s, \quad z_t = H_t W_t = H_s R W_t.$$

Unless $W_s = R W_t$, the two models can have identical CKA but different logits and different answer-token distributions. This distinction matches the difference between representational similarity and functional similarity discussed in prior work [56]. In heterogeneous VLM distillation, the final answer distribution also depends on the language head, tokenizer, output vocabulary and task-specific teacher behavior. Although SHM-CKA is competitive on ChartQA, it scores below GRACE on TextVQA and DocVQA, where output-level tokenizer-agnostic matching and sample-dependent teacher weighting are more important.

5.2 Representation Similarity Analysis

We use Centered Kernel Alignment (CKA) as an analysis tool to examine whether teacher and student layers exhibit clear correspondence. For centered hidden-state matrices $\tilde{\mathbf{H}}_s$ and $\tilde{\mathbf{H}}_t$, linear CKA is

$$\text{CKA}(\tilde{\mathbf{H}}_s, \tilde{\mathbf{H}}_t) = \frac{\left\| \tilde{\mathbf{H}}_t^\top \tilde{\mathbf{H}}_s \right\|_F^2}{\left\| \tilde{\mathbf{H}}_t^\top \tilde{\mathbf{H}}_t \right\|_F \left\| \tilde{\mathbf{H}}_s^\top \tilde{\mathbf{H}}_s \right\|_F}. \quad (5.1)$$

Values closer to 1 indicate higher representation similarity [19]. CKA is therefore used here as a representation-analysis tool, not as evidence that two models implement the same answer function. A high CKA score indicates similar hidden geometry on the evaluated samples, but it does not guarantee similar answer-token probabilities after the language head. The main question is whether high internal similarity is enough to justify layer-wise distillation. The results show that it is not, as the figures and tables below confirm.






Figure 5.6: Final-layer CKA heatmaps.

Figure 5.6 compares late hidden layers of Qwen2-VL-2B-Instruct and SmolVLM-500M. The TextVQA heatmap contains moderate similarities, mostly around 0.55 to 0.74, with higher values toward later student layers. The DocVQA heatmap is much brighter, with most entries around 0.87 to 0.93. These values show that the final language-side representations are not severely misaligned, especially on document-style inputs.



(a) TextVQA



(b) DocVQA

Figure 5.7: Vision-layer CKA heatmaps.

Figure 5.7 shows the same pattern in the vision stack. Both datasets contain broad high-similarity regions and the DocVQA vision map has higher CKA values across the middle layers. Teacher and student can therefore remain relatively close in both language and vision representations. However, the main benchmark results in Table 5.3 still contain a clear gap between the student models and higher-scoring teachers. The layer-wise baseline also fails to consistently outperform output-level distillation. This pattern supports the design choice in GRACE: output-level tokenizer-agnostic KD and sample-dependent teacher weighting produce higher benchmark scores than representation matching alone, because high hidden-state similarity does not guarantee that direct layer matching will produce better task transfer.

5.3 Ablation Study

This section answers RQ2 from Chapter 1 by isolating the main components of GRACE on DocVQA. Section 4.6 defines the ablation variants, with the corresponding results for both student models reported in Table 5.6.

Table 5.6: DocVQA component ablation.

DocVQA ablation variant	SmolVLM	
	256M	500M
Cross-family ULD (single teacher)	52.24	65.28
Uniform averaging	55.70	69.59
Routing only + trie-Wasserstein	57.45	63.23
Full GRACE + trie-Wasserstein	60.48	73.11
w/o EMA ($\beta_{\text{EMA}} = 0$)	57.45	71.25
w/ fixed blend ($\lambda_{\text{blend}} = 1.0$)	57.97	70.45

The ablation table should be read as a sequence of increasingly specific controls. The single-teacher ULD row tests whether cross-family token distribution transfer is useful before any teacher mixture is introduced. Uniform averaging then tests the value of teacher diversity without any adaptive weighting. The routing-only variant adds sample-conditioned teacher selection, while the full method adds the post-routing gradient check. The two final rows isolate the refinement module itself by removing temporal smoothing or by fixing the blend to rely only on routed weights over the active set. This ordering separates the benefit of using more teachers from the benefit of deciding which teacher signals should actually shape the update.

Teacher diversity is necessary but not sufficient because uniform multi-teacher averaging adds useful supervision while still failing to resolve tokenizer mismatch, teacher specialization and conflicting KD directions. Across the ablation results in Table 5.6, the full method is the highest-scoring setting for both student sizes. Moving from a single cross-family teacher (ULD) to uniform four-teacher averaging adds +3.46 on SmolVLM-256M and +4.31 on SmolVLM-500M. The gain shows that multiple heterogeneous teachers provide higher-scoring supervision than any single teacher, but uniform weighting leaves a further gap of +4.78/+3.52 to the full method.

Routing without gradient refinement can hurt. On SmolVLM-500M, routing-only scoring (63.23) is *worse* than uniform averaging (69.59) by -6.36 points. The router selects teachers but assigns them incorrect relative weights. Without GRACE, these wrong weights degrade the result below a naive equal average. On SmolVLM-256M the pattern reverses: routing-only (57.45) beats uniform (55.70) by $+1.75$. SmolVLM-500M, the larger student, produces a sharper routing distribution that amplifies weighting errors when uncorrected. SmolVLM-256M produces a flatter routing signal closer to uniform, so routing-only is less damaging and even slightly beneficial. This capacity-dependent behavior shows that a larger student representation does not automatically produce safer teacher weighting. A larger student may give the router a more separable representation of the sample, but the resulting gate can still be confidently wrong when teacher relevance and objective compatibility diverge. The smaller student appears to produce a less peaked routing signal, which behaves closer to a softened uniform mixture. This supports the design choice of treating routing as a proposal step and gradient-agreement refinement as a validation step.

Gradient-agreement refinement acts as the critical corrective layer because it adjusts routed teacher weights according to whether each teacher’s KD direction agrees with the supervised task objective. Full GRACE recovers $+9.88$ points over routing-only on SmolVLM-500M and $+3.03$ on SmolVLM-256M. Decomposing the 500M gain, the active-set router-weight variant ($\lambda_{\text{blend}} = 1.0$) already recovers $+7.22$ points ($63.23 \rightarrow 70.45$) by replacing uniform post-routing weights with learned router weights while retaining the GRACE active-set filter. Full score-softmax blending then adds another $+2.66$ points ($70.45 \rightarrow 73.11$). On SmolVLM-256M, the active-set router-weight variant adds only $+0.52$ points ($57.45 \rightarrow 57.97$), while full score-based blending adds $+2.51$ points ($57.97 \rightarrow 60.48$). These results show that active-set filtering and router weighting help, but agreement-score weighting is still needed for the highest final score. Removing EMA smoothing ($\beta_{\text{EMA}} = 0$) degrades both students (256M: $60.48 \rightarrow 57.45$, 500M: $73.11 \rightarrow 71.25$), confirming that temporal smoothing stabilizes score estimates but is secondary to the refinement step itself. The fixed-blend and no-EMA variants separate two roles of the refinement module. The fixed-blend variant keeps the active-set restriction but removes the score-softmax influence, so it tests whether filtering alone is enough. Its lower score shows that selecting an active teacher set is useful but incomplete. The no-EMA variant keeps score weighting but removes temporal smoothing, showing that batch-level gradient agreement is noisy and benefits from accumulated teacher-level history. Together, these two controls indicate that GRACE needs both a teacher-compatibility signal and a stable way to use that signal across training updates.

Answer to RQ2

The ablation isolates three cumulative contributions: (1) heterogeneous teacher diversity (+3 to +4 pts over single-teacher ULD), (2) sparse routing structure (beneficial only when paired with refinement) and (3) gradient-agreement refinement, which adds +3 to +10 pts over routing-only and is the single largest contributor. EMA smoothing provides additional stabilization by reducing the effect of noisy batch-level gradient-agreement scores before they are converted into final KD weights. Without gradient refinement, routing can degrade performance below uniform averaging.

5.4 Efficiency Analysis

This section addresses RQ3 from Chapter 1. Runtime, training cost and storage overhead are analyzed using the H200 rental rate and the observed Cloudflare R2 storage footprint. **Component-level overhead.** GRACE adds online overhead from three sources: sparse trie-Wasserstein support construction, teacher routing and gradient-agreement scoring. The trie-Wasserstein term is the dominant online algorithmic addition because it must compare sparse student and teacher output supports at supervised answer positions. Its cost depends on the retained top- K_{TW} support, the number of active trie edges and the number of routed teachers. The teacher router is comparatively small because it is an MLP over pooled student hidden states, so its cost is minor relative to the student forward and backward pass. Gradient-agreement refinement is more expensive than ordinary fixed weighting because it scores teacher-specific KD directions. This cost scales with the number of teacher-specific KD losses being scored, while top- k routing keeps the final KD aggregation sparse.

Offline versus online cost. The implementation separates teacher inference from student training. Teacher logits are generated offline and stored before distillation, so the online training loop does not need to run the frozen teachers during every student update. This reduces online GPU memory pressure and makes the reported experiments feasible on a single high-memory GPU. The tradeoff is that cost moves to the offline cache-generation stage and to storage and I/O. This design is attractive when the same cached teacher logits are reused across multiple ablations, since the expensive teacher forward passes are paid once. It becomes less attractive when datasets, prompts or teacher pools change frequently, because the corresponding cached logits must be regenerated.

Storage and cache tradeoff. Figure B.1 shows the R2 bucket used for teacher-logit caches. The full store occupies 1.04 TB, which is too large to copy repeatedly to temporary external GPU instances. The training pipeline therefore streams logits through the local cache described in Appendix B.3. This makes the local disk requirement bounded by the 100 GB cache limit instead of by the full remote store. The cost is additional I/O latency on cache misses and a dependence on remote object storage availability. In practice, this tradeoff is acceptable for the reported runs because storage cost is small compared with GPU rental cost.

Scalability. The main scaling pressure comes from the product of dataset size, number of teachers and cached answer-token logits. Adding teachers increases offline generation time, remote storage and the number of candidate KD losses that may need to be scored. Increasing K_{TW} raises the online cost of sparse trie construction, while increasing k_{route} raises the number of teacher losses used per sample. The current four-teacher setting is a workable operating point in this tradeoff: it is large enough to expose teacher specialization, but still small enough for sparse routing, gradient refinement and remote-cache streaming to remain manageable.

Answer to RQ3

GRACE introduces online overhead through trie-Wasserstein sparse matching, teacher routing and gradient-agreement scoring, with trie matching and teacher-specific gradient scoring being the main additions beyond ordinary fine-tuning. The implementation avoids running frozen teachers during student training by using offline teacher-logit caches. This shifts part of the cost to offline logit generation and remote storage. The observed R2 store occupies 1.04 TB and costs about 19.31 USD/month before rounding, while a 7-hour H200 run costs about 14.41 USD. The dominant bottleneck is therefore GPU compute and offline logit generation, not object storage.

5.5 Discussion

The score tables show the same ordering across method families. Same-family KL baselines obtain the lowest scores, single cross-family distillation scores higher, uniform multi-teacher weighting improves further and GRACE obtains the highest scores overall [11, 16]. No single component fully addresses the heterogeneous distillation problem: heterogeneous teachers improve scores, but the student still needs a tokenizer-agnostic loss, sample-dependent routing and a stable post-routing weighting scheme in the evaluated setting. The training recipe explains this pattern because GRACE combines a shared-trie Wasserstein loss for mismatched tokenizers with routed multi-teacher selection, top- k masking, routed-weight renormalization and EMA-smoothed gradient-agreement refinement [16, 32, 36]. The smoother DocVQA training curve in Figure 5.4 is consistent with that design: the router narrows the teacher set first, then GRACE reweights only the routed teachers, reducing the score oscillations caused by uniformly mixing heterogeneous teacher signals. **First-order interpretation of gradient-agreement refinement.** The gradient-agreement component can also be understood from a first-order multi-objective view. Let $g_{CE} = \nabla_{\theta} \mathcal{L}_{CE}$ and $g_i = \nabla_{\theta} \mathcal{L}_{KD}^{(i)}$ be the supervised and teacher-specific KD gradients. A combined update with teacher weights w_i is

$$\Delta\theta = -\eta \left(g_{CE} + \alpha_{KD} \sum_i w_i g_i \right).$$

The first-order change in the supervised loss is

$$\Delta\mathcal{L}_{CE} \approx g_{CE}^{\top} \Delta\theta = -\eta \|g_{CE}\|_2^2 - \eta \alpha_{KD} \sum_i w_i \langle g_{CE}, g_i \rangle.$$

If $\langle g_{CE}, g_i \rangle < 0$, then teacher i contributes a term that increases the first-order CE loss when the supervised objective should decrease. This is the gradient-conflict case studied in multi-task optimization [41]. GRACE does not project gradients as PCGrad does. It uses the cosine agreement

$$\cos(g_{CE}, g_i) = \frac{\langle g_{CE}, g_i \rangle}{\|g_{CE}\|_2 \|g_i\|_2}$$

as a teacher-weighting signal. This explains why post-routing gradient refinement can improve over routing alone. The largest gain appears on ChartQA and the trend holds for SmolVLM-256M as well. Sample-dependent teacher assignment has the largest observed benefit when the task mixes visual parsing and numerical reasoning. The method also benefits both evaluated student models. The comparison with SHM-CKA reinforces this conclusion by showing that heterogeneous teacher access is not enough when the method relies only on hidden-state matching with no routed output-level KD. SHM-CKA is competitive on ChartQA yet scores lower than GRACE on TextVQA and DocVQA. Together with the CKA heatmaps, this confirms that internal similarity alone does not guarantee effective transfer. **Sample-dependent teacher weighting therefore remains necessary, particularly when the teacher pool contains models with different tokenizers, architectures and task-specific strengths** [19, 31]. The remaining limitations and future directions are discussed in the concluding chapter.

Conclusion and Future Work

Summary

This thesis studied knowledge distillation for heterogeneous vision-language models, where teacher utility varies across samples and teacher-student tokenizers do not match. The work introduced Gradient-Routed Alignment with Capacity-aware Experts (GRACE), a routed multi-teacher framework that combines three main contributions: trie-Wasserstein tokenizer-agnostic distillation, sparse teacher routing and gradient-agreement refinement [16, 32, 36].

The first contribution is a geometry-aware cross-tokenizer distillation objective. Trie-Wasserstein KD maps decoded teacher and student tokens into a shared byte-trie structure and compares subtree probability masses, so token distributions can be matched across different vocabularies. The second contribution is an objective-aware post-routing refinement mechanism. Sparse routing selects sample-relevant teachers, while gradient-agreement refinement adjusts routed teacher weights according to alignment between each teacher’s KD gradient and the supervised answer gradient. The third contribution is an empirical study showing that output-level tokenizer-agnostic distillation with routing and gradient agreement outperforms the compared same-family, single-teacher and heterogeneous multi-teacher baselines in the evaluated VQA setting.

Across TextVQA, DocVQA and ChartQA, GRACE obtains the highest scores among the compared methods in Chapter 5. For SmolVLM-500M, it reaches 69.38 on TextVQA, 73.11 on DocVQA and 80.20 on ChartQA. For SmolVLM-256M, it reaches 60.73, 60.48 and 72.32 on the same three benchmarks. On SmolVLM-500M, GRACE closes about 55% of the original teacher-student gap on TextVQA and ChartQA and about 40% on DocVQA.

Limitations

Several aspects remain unresolved. **Cost-limited evaluation** is the first limitation. The reported numbers are single-run results, so the thesis does not provide confidence intervals or repeated-seed estimates for small gaps between high-scoring baselines. As shown in Section 5.4, an average run costs approximately 14.41 USD before offline cache generation and storage cost. Although the main tables use a shared four-teacher pool and consistent hyperparameters, the absence of repeated seeds limits strict uncertainty claims about closely matched methods.

Scope of validation is also limited. The empirical study focuses on VQA benchmarks and SmolVLM-family students. The results support text-rich VQA distillation with compact projector-based students, but they do not establish whether the same behavior holds for other multimodal task families, other student architectures or larger teacher pools. Tasks with long-form answers, multi-image inputs, video reasoning or open-ended dialogue may place different demands on teacher selection and answer-token supervision.

Reasoning-oriented models are not explicitly supported. GRACE currently distills teacher supervision through answer-token output distributions and gradient agreement with the supervised answer loss. This design is appropriate for the evaluated VQA setting, where the target is a short final answer. It does not yet explicitly support reasoning-oriented or “thinking” models whose behavior depends on intermediate deliberation traces, hidden reasoning tokens or extended test-time reasoning. If a teacher produces only a final answer distribution, GRACE can still use that distribution as a teacher signal. However, the method does not distinguish between reasoning-phase tokens and answer-phase tokens, does not route teachers based on the quality of intermediate reasoning and does not define a distillation objective for multi-step thought trajectories.

Future Work

Future work should extend the thesis in four directions. First, stronger evaluation should use repeated runs with multiple random seeds, report mean and standard deviation and use paired comparisons when the same cached teacher logits are reused. This would make the smaller differences between close multi-teacher baselines easier to interpret.

Second, systems work should reduce the storage, I/O and gradient-computation costs of heterogeneous teacher distillation. More compact logit storage, adaptive cache loading and lower-cost agreement scoring would make the method easier to scale to more datasets and larger teacher pools.

Third, future method extensions should move beyond answer-token-centric distillation. Visual-token alignment, cross-attention transfer and region-level matching could complement output-level KD, especially for tasks where the student must localize evidence before producing a short answer.

Fourth, support for reasoning-oriented models should be studied with process-level distillation mechanisms. Such extensions would need to separate reasoning-stage information from final-answer generation and define how teachers are routed or weighted when intermediate reasoning traces are available. Future work can also explore richer router inputs, improved objective-conflict handling and teacher pools that mix specialized document, chart, general VLM and reasoning-oriented experts [20–22].

References

- [1] Jean-Baptiste Alayrac, Jeff Donahue, Pauline Luc, Antoine Miech, Iain Barr, Yana Hasson, Karel Lenc, Arthur Mensch, Katie Millican, Malcolm Reynolds, Roman Ring, Eliza Rutherford, Serkan Cabi, Tengda Han, Zhitao Gong, Sina Samangooei, Marianne Monteiro, Jacob Menick, Sebastian Borgeaud, Andrew Brock, Aida Nematzadeh, Sahand Sharifzadeh, Mikolaj Binkowski, Ricardo Barreira, Oriol Vinyals, Andrew Zisserman, and Karen Simonyan. Flamingo: a visual language model for few-shot learning, 2022.
- [2] Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. Blip-2: Bootstrapping language-image pre-training with frozen image encoders and large language models, 2023.
- [3] Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning, 2023.
- [4] Amanpreet Singh, Vivek Natarajan, Meet Shah, Yu Jiang, Xinlei Chen, Dhruv Batra, Devi Parikh, and Marcus Rohrbach. Towards vqa models that can read, 2019.
- [5] Minesh Mathew, Dimosthenis Karatzas, and C. V. Jawahar. Docvqa: A dataset for vqa on document images, 2021.
- [6] Ahmed Masry, Do Xuan Long, Jia Qing Tan, Shafiq Joty, and Enamul Hoque. Chartqa: A benchmark for question answering about charts with visual and logical reasoning, 2022.
- [7] Danny Driess, Fei Xia, Mehdi S. M. Sajjadi, Corey Lynch, Aakanksha Chowdhery, Brian Ichter, Ayzaan Wahid, Jonathan Tompson, Quan Vuong, Tianhe Yu, Wenlong Huang, Yevgen Chebotar, Pierre Sermanet, Daniel Duckworth, Sergey Levine, Vincent Vanhoucke, Karol Hausman, Marc Toussaint, Klaus Greff, Andy Zeng, Igor Mordatch, and Pete Florence. Palm-e: An embodied multimodal language model, 2023.

- [8] Anthony Brohan, Noah Brown, Justice Carbajal, Yevgen Chebotar, Xi Chen, Krzysztof Choromanski, Tianli Ding, Danny Driess, Avinava Dubey, Chelsea Finn, Pete Florence, Chuyuan Fu, Montse Gonzalez Arenas, Keerthana Gopalakrishnan, Kehang Han, Karol Hausman, Alexander Herzog, Jasmine Hsu, Brian Ichter, Alex Irpan, Nikhil Joshi, Ryan Julian, Dmitry Kalashnikov, Yuheng Kuang, Isabel Leal, Lisa Lee, Tsang-Wei Edward Lee, Sergey Levine, Yao Lu, Henryk Michalewski, Igor Mordatch, Karl Pertsch, Kanishka Rao, Krista Reymann, Michael Ryoo, Grecia Salazar, Pannag Sanketi, Pierre Sermanet, Jaspiar Singh, Anikait Singh, Radu Soricut, Huong Tran, Vincent Vanhoucke, Quan Vuong, Ayzaan Wahid, Stefan Welker, Paul Wohlhart, Jialin Wu, Fei Xia, Ted Xiao, Peng Xu, Sichun Xu, Tianhe Yu, and Brianna Zitkovich. Rt-2: Vision-language-action models transfer web knowledge to robotic control, 2023.
- [9] Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. Scaling laws for neural language models, 2020.
- [10] Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, Tom Hennigan, Eric Noland, Katie Millican, George van den Driessche, Bogdan Damoc, Aurelia Guy, Simon Osindero, Karen Simonyan, Erich Elsen, Jack W. Rae, Oriol Vinyals, and Laurent Sifre. Training compute-optimal large language models, 2022.
- [11] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network, 2015.
- [12] Xiaohan Xu, Ming Li, Chongyang Tao, Tao Shen, Reynold Cheng, Jinyang Li, Can Xu, Dacheng Tao, and Tianyi Zhou. A survey on knowledge distillation of large language models, 2024.
- [13] Xiangxiang Chu, Limeng Qiao, Xinyang Lin, Shuang Xu, Yang Yang, Yiming Hu, Fei Wei, Xinyu Zhang, Bo Zhang, Xiaolin Wei, and Chunhua Shen. Mobilevlm: A fast, strong and open vision language assistant for mobile devices, 2023.
- [14] Baichuan Zhou, Ying Hu, Xi Weng, Junlong Jia, Jie Luo, Xien Liu, Ji Wu, and Lei Huang. Tinyllava: A framework of small-scale large multimodal models, 2024.

- [15] Andrés Marafioti, Orr Zohar, Miquel Farré, Merve Noyan, Elie Bakouch, Pedro Cuenca, Cyril Zakka, Loubna Ben Allal, Anton Lozhkov, Nouamane Tazi, Vaibhav Srivastav, Joshua Lochner, Hugo Larcher, Mathieu Morlon, Lewis Tunstall, Leandro von Werra, and Thomas Wolf. Smolvlm: Redefining small and efficient multimodal models, 2025.
- [16] Nicolas Boizard, Kevin El Haddad, Pierre Colombo, and Chloé Clavel. Towards cross-tokenizer distillation: The universal logit distillation loss for llms, 2024.
- [17] Yuki Takezawa, Ryoma Sato, and Makoto Yamada. Supervised tree-wasserstein distance. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, page 10086–10095. PMLR, 2021.
- [18] Tuan Le, Makoto Yamada, Kenji Fukumizu, and Marco Cuturi. Tree-sliced variants of wasserstein distances. In *Advances in Neural Information Processing Systems*, volume 32, 2019.
- [19] Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network representations revisited. In *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, page 3519–3529. PMLR, 2019.
- [20] Fei Yuan, Linjun Shou, Jian Pei, Wutao Lin, Ming Gong, Yan Fu, and Daxin Jiang. Reinforced multi-teacher selection for knowledge distillation, 2020.
- [21] Hailin Zhang, Defang Chen, and Can Wang. Adaptive multi-teacher knowledge distillation with meta-learning, 2023.
- [22] Yu-Chu Yu, Chi-Pin Huang, Jr-Jen Chen, Kai-Po Chang, Yung-Hsuan Lai, Fu-En Yang, and Yu-Chiang Frank Wang. Select and distill: Selective dual-teacher knowledge transfer for continual learning on vision-language models, 2024.
- [23] Peng Wang, Shuai Bai, Sinan Tan, Shijie Wang, Zhihao Fan, Jinze Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Yang Fan, Kai Dang, Mengfei Du, Xuancheng Ren, Rui Men, Dayiheng Liu, Chang Zhou, Jingren Zhou, and Junyang Lin. Qwen2-vl: Enhancing vision-language model’s perception of the world at any resolution, 2024.

- [24] Zhe Chen, Weiyun Wang, Hao Tian, Shenglong Ye, Zhangwei Gao, Erfei Cui, Wenwen Tong, Kongzhi Hu, Jiapeng Luo, Zheng Ma, Ji Ma, Jiaqi Wang, Xiaoyi Dong, Hang Yan, Hewei Guo, Conghui He, Botian Shi, Zhenjiang Jin, Chao Xu, Bin Wang, Xingjian Wei, Wei Li, Wenjian Zhang, Bo Zhang, Pinlong Cai, Licheng Wen, Xiangchao Yan, Min Dou, Lewei Lu, Xizhou Zhu, Tong Lu, Dahua Lin, Yu Qiao, Jifeng Dai, and Wenhai Wang. How far are we to gpt-4v? closing the gap to commercial multimodal models with open-source suites, 2024.
- [25] Zhe Chen, Weiyun Wang, Yue Cao, Yangzhou Liu, Zhangwei Gao, Erfei Cui, Jinguo Zhu, Shenglong Ye, Hao Tian, Zhaoyang Liu, Lixin Gu, Xuehui Wang, Qingyun Li, Yiming Ren, Zixuan Chen, Jiapeng Luo, Jiahao Wang, Tan Jiang, Bo Wang, Conghui He, Botian Shi, Xingcheng Zhang, Han Lv, Yi Wang, Wenqi Shao, Pei Chu, Zhongying Tu, Tong He, Zhiyong Wu, Huipeng Deng, Jiaye Ge, Kai Chen, Kaipeng Zhang, Limin Wang, Min Dou, Lewei Lu, Xizhou Zhu, Tong Lu, Dahua Lin, Yu Qiao, Jifeng Dai, and Wenhai Wang. Expanding performance boundaries of open-source multimodal models with model, data, and test-time scaling, 2024.
- [26] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision, 2021.
- [27] Xiaohua Zhai, Basil Mustafa, Alexander Kolesnikov, and Lucas Beyer. Sigmoid loss for language image pre-training. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, page 11975–11986, 2023.
- [28] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems 30 (NeurIPS 2017)*, page 5998–6008, 2017.
- [29] Jianhua Lin. Divergence measures based on the shannon entropy. *IEEE Transactions on Information Theory*, 37(1), 1991.
- [30] Zhiyuan Fang, Jianfeng Wang, Xiaowei Hu, Lijuan Wang, Yezhou Yang, and Zicheng Liu. Compressing visual-linguistic model via knowledge distillation. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, page 1428–1438, 2021.

- [31] Sayantan Dasgupta and Trevor Cohn. Improving language model distillation through hidden state matching. In *International Conference on Learning Representations*, 2025.
- [32] Gabriel Peyré and Marco Cuturi. Computational optimal transport, 2019.
- [33] Rémi Flamary and Nicolas Courty. Optimal transport for machine learning: Introduction to optimal transport, 2019. Tutorial at ISBI 2019.
- [34] Wikimedia Commons contributors. File: Optimal transport matrix.png. Wikimedia Commons, 2018. CC BY-SA 4.0. Accessed: 2026-03-31.
- [35] Noam Shazeer, Azalia Mirhoseini, Krzysztof Maziarczyk, Andy Davis, Quoc Le, Geoffrey Hinton, and Jeff Dean. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer, 2017.
- [36] William Fedus, Barret Zoph, and Noam Shazeer. Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity, 2022.
- [37] DeepSeek-AI. Deepseek-v2: A strong, economical and efficient mixture-of-experts language model, 2024.
- [38] Dmitry Lepikhin, Hyoungho Lee, Yuanzhong Xu, Dehao Chen, Orhan Firat, Yanping Huang, Maxim Krikun, Noam Shazeer, and Zhifeng Chen. Gshard: Scaling giant models with conditional computation and automatic sharding, 2020.
- [39] Yanqi Zhou, Tao Lei, Hanxiao Liu, Nan Du, Yanping Huang, Vincent Zhao, Andrew Dai, Zhifeng Chen, Quoc Le, and James Laudon. Mixture-of-experts with expert choice routing, 2022.
- [40] Yichen Zhu and Yi Wang. Student customized knowledge distillation: Bridging the gap between student and teacher. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, page 5057–5066, 2021.
- [41] Tianhe Yu, Saurabh Kumar, Abhishek Gupta, Sergey Levine, Karol Hausman, and Chelsea Finn. Gradient surgery for multi-task learning. In *Advances in Neural Information Processing Systems*, volume 33, page 5824–5836, 2020.
- [42] Yushi Hu, Otilia Stretcu, Chun-Ta Lu, Krishnamurthy Viswanathan, Kenji Hata, Enming Luo, Ranjay Krishna, and Ariel Fuxman. Visual program distillation: Distilling tools and programmatic reasoning into vision-language models, 2024.

- [43] Kaicheng Yang, Tiancheng Gu, Xiang An, Haiqiang Jiang, Xiangzi Dai, Ziyong Feng, Weidong Cai, and Jiankang Deng. Clip-cid: Efficient clip distillation via cluster-instance discrimination, 2024.
- [44] Haodong Duan, Xinyu Fang, Junming Yang, et al. Vlmevalkit: An open-source toolkit for evaluating large multi-modality models, 2024.
- [45] Wikimedia Commons contributors. File: Stop sign.jpg. Wikimedia Commons, 2006. Public domain. Accessed: 2026-03-06.
- [46] Wikimedia Commons contributors. File: Receiptswiss.jpg. Wikimedia Commons, 2015. CC BY-SA 3.0. Accessed: 2026-03-06.
- [47] Wikimedia Commons contributors. File: Bar chart.svg. Wikimedia Commons, 2016. CC BY-SA 4.0. Accessed: 2026-03-06.
- [48] Samyam Rajbhandari, Jeff Rasley, Olatunji Ruwase, and Yuxiong He. Zero: Memory optimizations toward training trillion parameter models, 2020.
- [49] Tianqi Chen, Bing Xu, Chiyuan Zhang, and Carlos Guestrin. Training deep nets with sublinear memory cost, 2016.
- [50] Hugging Face TB. Smolvlm2-2.2b-instruct model card. Hugging Face, 2025. Accessed: 2026-04-02.
- [51] Yifan Cui, Wei Wang, Yifei Wang, Yao Lu, Haoyu Wang, Fuzhen Zhuang, Qing He, and Chuan Shi. Sinkhorn distance minimization for knowledge distillation. In *Findings of the Association for Computational Linguistics: ACL 2024*, page 11046–11060. Association for Computational Linguistics, 2024.
- [52] Ronald J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine Learning*, 8(3):229–256, 1992.
- [53] Yuqiao Wen, Zichao Li, Wenyu Du, and Lili Mou. f-divergence minimization for sequence-level knowledge distillation. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, page 10817–10834. Association for Computational Linguistics, 2023.
- [54] Walid Bousselham, Rishikesh Jadhav, Simon de Mathelin, François Fleuret, and Pascal Frossard. Dex-ar: Explainability in multimodal large language models with deep expectation and attention rollout, 2026.

- [55] Taiqiang Wu, Chaofan Tao, Jiahao Wang, Runming Yang, Zhe Zhao, and Ngai Wong. Rethinking kullback-leibler divergence in knowledge distillation for large language models. In *Proceedings of the 31st International Conference on Computational Linguistics*, page 5737–5755, 2025.
- [56] Max Klabunde, Tobias Schumacher, Markus Strohmaier, and Florian Lemmerich. Similarity of neural network models: A survey of functional and representational measures, 2023.
- [57] Cloudflare. R2 object storage. Cloudflare developer documentation, 2026. Accessed: 2026-04-09.
- [58] Cloudflare. Use the aws sdk for python (boto3) with r2. Cloudflare developer documentation, 2026. Accessed: 2026-04-09.
- [59] fsspec developers. Features. fsspec documentation, 2026. Accessed: 2026-04-09.
- [60] fsspec developers. fsspec.implementations.cached. fsspec documentation, 2026. Accessed: 2026-04-09.

Appendix A

Proof of the Closed-Form Trie-Wasserstein

This appendix derives the closed-form expression used by the KD loss in Chapter 3. The result is stated for a general rooted tree [17, 32].

Setup.

Let $T = (V, E, r)$ be a rooted tree with root r . Each edge $e \in E$ has a positive weight $w(e) > 0$. For an edge $e = (p(e), c(e))$, where $p(e)$ is the parent and $c(e)$ is the child, let

$$\Gamma(e) = \{v \in V \mid e \text{ lies on the path from } r \text{ to } v\}$$

denote the subtree below edge e . For two vertices $u, v \in V$, define the tree metric

$$d_T(u, v) = \sum_{e \in \text{path}(u, v)} w(e),$$

where $\text{path}(u, v)$ is the unique set of edges on the path between u and v . Let μ and ν be two probability measures on V . In the trie-Wasserstein setting, these correspond to student and teacher token distributions placed on the leaves of the shared byte trie. Measures on leaves can be viewed as measures on all vertices by assigning zero mass to non-leaf vertices. The Kantorovich form of the W_1 optimal transport distance under the tree metric is [32]

$$W_{1,T}(\mu, \nu) = \min_{\pi \in \Pi(\mu, \nu)} \sum_{u \in V} \sum_{v \in V} \pi(u, v) d_T(u, v),$$

where $\Pi(\mu, \nu)$ is the set of couplings satisfying

$$\sum_{v \in V} \pi(u, v) = \mu(u), \quad \sum_{u \in V} \pi(u, v) = \nu(v), \quad \pi(u, v) \geq 0.$$

For each edge e , define the subtree-mass imbalance

$$\Delta(e) = \mu(\Gamma(e)) - \nu(\Gamma(e)).$$

Hence, the closed-form expression for the tree Wasserstein distance is

$$W_{1,T}(\mu, \nu) = \sum_{e \in E} w(e) |\Delta(e)|.$$

Lower bound.

Consider any feasible transport plan $\pi \in \Pi(\mu, \nu)$. Removing an edge e separates the tree into two components: the subtree $\Gamma(e)$ and its complement $V \setminus \Gamma(e)$. The total source mass initially inside $\Gamma(e)$ is $\mu(\Gamma(e))$, while the total target mass required inside $\Gamma(e)$ is $\nu(\Gamma(e))$. If

$$\mu(\Gamma(e)) > \nu(\Gamma(e)),$$

then at least $\mu(\Gamma(e)) - \nu(\Gamma(e))$ mass must leave $\Gamma(e)$. If

$$\mu(\Gamma(e)) < \nu(\Gamma(e)),$$

then at least $\nu(\Gamma(e)) - \mu(\Gamma(e))$ mass must enter $\Gamma(e)$. Therefore, any feasible plan must move at least

$$|\Delta(e)| = |\mu(\Gamma(e)) - \nu(\Gamma(e))|$$

mass across edge e . Let $q_\pi(e)$ be the amount of mass transported by π whose path crosses edge e . The argument above implies

$$q_\pi(e) \geq |\Delta(e)|.$$

Because the cost of moving mass from u to v is the sum of edge weights along the unique path between them, the total cost of π can be decomposed edgewise:

$$\sum_{u,v \in V} \pi(u,v) d_T(u,v) = \sum_{e \in E} w(e) q_\pi(e).$$

Using $q_\pi(e) \geq |\Delta(e)|$, we obtain

$$\sum_{u,v \in V} \pi(u,v) d_T(u,v) \geq \sum_{e \in E} w(e) |\Delta(e)|.$$

Since this holds for every feasible transport plan π , it also holds for the optimal plan:

$$W_{1,T}(\mu, \nu) \geq \sum_{e \in E} w(e) |\Delta(e)|.$$

Achievability.

It remains to show that the lower bound can be achieved. We define a signed measure

$$\sigma(v) = \mu(v) - \nu(v).$$

For each edge e , the subtree imbalance is

$$\Delta(e) = \sum_{v \in \Gamma(e)} \sigma(v).$$

Interpret $\Delta(e)$ as the net amount of mass that must cross edge e . If $\Delta(e) > 0$, the subtree below e contains surplus mass and $\Delta(e)$ mass must move from $\Gamma(e)$ to its complement. If $\Delta(e) < 0$, the subtree has a deficit and $|\Delta(e)|$ mass must enter $\Gamma(e)$. Because T is a tree, there is a unique path between any two vertices. A bottom-up transport construction can move surplus mass upward from subtrees with $\Delta(e) > 0$ and move mass downward into subtrees with $\Delta(e) < 0$. This construction sends exactly $|\Delta(e)|$ mass across each edge e . It satisfies all supply and demand constraints because the total signed mass is zero:

$$\sum_{v \in V} \sigma(v) = \sum_{v \in V} \mu(v) - \sum_{v \in V} \nu(v) = 1 - 1 = 0.$$

The cost of this construction is therefore

$$\sum_{e \in E} w(e) |\Delta(e)|.$$

Thus there exists a feasible transport plan with cost equal to the lower bound. Hence

$$W_{1,T}(\mu, \nu) \leq \sum_{e \in E} w(e) |\Delta(e)|.$$

Combining this upper bound with the lower bound gives

$$W_{1,T}(\mu, \nu) = \sum_{e \in E} w(e) |\mu(\Gamma(e)) - \nu(\Gamma(e))|.$$

Appendix B

Training Configuration and Infrastructure

This appendix lists the optimization and routing hyperparameters used across all distillation experiments and records the engineering details needed to reproduce the training infrastructure. The hyperparameters correspond to the evaluation protocol described in Chapter 4.

Table B.1: Experiment parameters.

Setting	Value
<i>Model setup</i>	
Student model	HuggingFaceTB/SmolVLM-500M-Instruct
Teacher models	Qwen/Qwen2.5-VL-3B-Instruct, Qwen/Qwen2-VL-2B-Instruct, OpenGVLab/InternVL3-1B, google/gemma-3-4b-it
Teacher weighting strategy	routing (GRACE path)
Distillation loss	Trie-Wasserstein
Trie-Wasserstein rho	0.7
Trie-Wasserstein top- k	64
Validation split	validation
<i>Hardware and infrastructure</i>	
Hardware and runtime	Single-GPU H200 or Vast.ai runs
Distributed training	DeepSpeed ZeRO-2
Numerical precision	bf16 enabled, half precision (fp16) disabled, TF32 enabled
FlashAttention-2	Enabled
Training data format	LLaVA-style multimodal instruction JSON
<i>Training</i>	
Max sequence length	16384
Training epochs	2
Per-device batch size	80
Gradient accumulation steps	1
Effective batch size	80
Gradient checkpointing	Enabled
Frozen modules	None
Dataloader workers	4
Logging and checkpointing	Log every 5 steps and save every 150 steps

Continued on the next page.

Table B.1 continued from the previous page.

Setting	Value
Seed	42
<i>Optimizer</i>	
Optimizer	adamw_torch
Adam betas	(0.9, 0.98)
Adam epsilon	1×10^{-7}
Weight decay	0.0
Base learning rate	1×10^{-5}
Vision learning rate	2×10^{-6}
Connector learning rate	1×10^{-5}
Learning-rate schedule	Cosine decay with warmup ratio 0.03
<i>GRACE hyperparameters</i>	
Student/teacher temperature	$\tau_s = \tau_t = 1.0$
KD weight	$\alpha_{\text{KD}} = 0.5$
Top- k routing	$k_{\text{route}} = 2$, with top-scoring teacher masking and selected softmax-weight renormalization
Gate entropy weight	$\lambda_{\text{ent}} = 10^{-3}$
Router z -loss weight	$\lambda_z = 10^{-3}$
GRACE warmup ratio	$\rho_g = 0.05$
GRACE softmax beta	$\beta_{\text{score}} = 4.0$
GRACE blend exponent	$\lambda_{\text{blend}} = 0.9$
GRACE EMA decay	$\beta_{\text{EMA}} = 0.8$
Score-spread fallback threshold	$\epsilon_{\text{fb}} = 0.01$
Alignment threshold	$\delta = 0.15$

B.1 Baseline Implementation

This section records the additional baseline definitions used for the comparison methods in Chapter 4, including how their teacher selection, weighting and auxiliary loss terms are implemented for fair comparison.

SHM-CKA. The baseline first loads an offline CKA matrix and, for each student layer s , selects the top- k teacher layers $\mathcal{T}_s$. Their CKA scores are normalized into soft weights $\alpha_{s,t}$. The per-teacher layer loss is

$$\mathcal{L}_{\text{layer}}^{(i)} = \frac{1}{|S|} \sum_{s \in S} \sum_{t \in \mathcal{T}_s} \alpha_{s,t} (1 - \sqrt{\text{CKA}(\mathbf{H}_s, \mathbf{H}_t)}). \quad (\text{B.1})$$

Each teacher also contributes a ULD logits loss $\mathcal{L}_{\text{ULD}}^{(i)}$. For the heterogeneous baseline, the logits and layer-matching terms are averaged across teachers:

$$\bar{\mathcal{L}}_{\text{ULD}} = \frac{1}{M} \sum_{i=1}^M \mathcal{L}_{\text{ULD}}^{(i)}, \quad \bar{\mathcal{L}}_{\text{layer}} = \frac{1}{M} \sum_{i=1}^M \mathcal{L}_{\text{layer}}^{(i)}. \quad (\text{B.2})$$

The final SHM-CKA objective is

$$\mathcal{L}_{\text{SHM-CKA}} = \frac{1}{3} (\mathcal{L}_{\text{CE}} + \bar{\mathcal{L}}_{\text{ULD}} + \bar{\mathcal{L}}_{\text{layer}}). \quad (\text{B.3})$$

REINFORCE. The policy uses the pooled student hidden state together with per-teacher summary features. A linear policy predicts independent Bernoulli probabilities

$$\boldsymbol{\pi} = \sigma(W \text{LN}(\mathbf{z}) + b), \quad (\text{B.4})$$

then samples a teacher mask $m_i \sim \text{Bernoulli}(\pi_i)$. If no teacher is selected, the highest-probability teacher is used as fallback. The selected mask is normalized into teacher weights

$$w_i = \frac{m_i}{\sum_{j=1}^M m_j}, \quad \mathcal{L}_{\text{KD}} = \sum_{i=1}^M w_i \mathcal{L}_{\text{KD}}^{(i)}. \quad (\text{B.5})$$

The selector is trained with the policy-gradient term

$$\mathcal{L}_{\text{pol}} = -(r - b) \log p(\mathbf{m} \mid \boldsymbol{\pi}), \quad (\text{B.6})$$

where r is the reward and b is an EMA reward baseline. The total reinforced-selection objective is

$$\mathcal{L}_{\text{RS}} = \mathcal{L}_{\text{CE}} + \gamma \mathcal{L}_{\text{KD}} + \lambda_{\pi} \mathcal{L}_{\text{pol}}. \quad (\text{B.7})$$

B.2 Offline Teacher-Logit Preparation

This section records the offline teacher-logit generation step used before distillation training. An average full H200 student-training run takes about 7 hours and costs about 14.41 USD at 2.059 USD/hour on Vast.ai, excluding the separate cost of offline teacher-logit generation and remote storage.

Default Storage Class ⓘ	Public Access ⓘ	Bucket Size	Class A Operations ⓘ	Class B Operations ⓘ
Standard	Disabled	1.04 TB	584.51k	17.42M

ObjectsMetricsSettings

Search objects by prefix

Q Search

View prefixes as directories ⓘ

[google-drive-backup](#) / 1W9sUXnqNdR2qVHr8PrgVAtM6J2SXsJ-7 /

Upload

Add directory

☐ Objects	Type	Storage Class	Size	Modified
☐ chartqa_teacher_logits_gemma3_4b/	Directory	--		
☐ chartqa_teacher_logits_intervl2_1b/	Directory	--		
☐ chartqa_teacher_logits_qwen25vl_3b/	Directory	--		
☐ chartqa_teacher_logits_qwen2vl_2b/	Directory	--		
☐ docvqa_teacher_logits_gemma3_4b/	Directory	--		
☐ docvqa_teacher_logits_granite_vision_3_1_2b_preview/	Directory	--		
☐ docvqa_teacher_logits_intervl2_1b/	Directory	--		
☐ docvqa_teacher_logits_qwen25vl_3b/	Directory	--		
☐ docvqa_teacher_logits_qwen2vl_2b/	Directory	--		
☐ streaming_teacher_logits_smoke/	Directory	--		
☐ textvqa_teacher_logits_gemma3_4b/	Directory	--		
☐ textvqa_teacher_logits_intervl2_1b/	Directory	--		
☐ textvqa_teacher_logits_qwen25vl_3b/	Directory	--		
☐ textvqa_teacher_logits_qwen2vl_2b/	Directory	--		

Figure B.1: R2 teacher-logit storage.

Teacher logits are generated *offline* before distillation. An offline caching script runs each frozen teacher on every training example, keeps only the supervised answer-token positions and writes one `.pt` file per sample under a teacher-specific directory. Each record stores the sample index, the selected teacher logits and the corresponding supervised labels. This keeps the online training loop simple: the student only reads precomputed teacher targets, so the teachers do not run again. After generation, the teacher-logit directories are uploaded to a **Cloudflare R2** bucket using its S3-compatible object-storage interface [57, 58]. In the active setup, the bucket already exceeds **1 TB**, so copying the full cache to an external GPU instance is not feasible. This is especially relevant on services such as Vast.ai, where local NVMe space is usually much smaller than the full teacher logit store. Figure B.1 shows this offline storage layout and highlights why the training path must stream teacher logits without staging the full cache locally.




Figure B.2: Training with remote teacher logits.

B.3 Remote Logit Streaming During Training

The main training path does *not* download the full cache before training. Figure B.2 shows the full remote-cache path: teacher logits are generated offline and uploaded once, each training step requests a sample, the local cache serves hits directly and cache misses trigger an R2 download followed by least-recently-used eviction when the 100 GB limit is exceeded. When the launcher receives an R2 root in `distill_multi_teachers.sh`, the trainer switches to remote lazy loading. During training, `StreamingTeacherLogitsCache` resolves the teacher-specific remote object path for each requested sample, then fetches files lazily from R2 through `s3fs` and `WholeFileCacheFilesystem` [59, 60]. To keep disk usage bounded, the loader maintains a **100 GB** local cache. Each file is cached on first access, its access time is updated on each subsequent open and the oldest files are evicted once the cache exceeds the limit. The local cache therefore behaves as a small least-recently-used buffer over the much larger R2 store. A per-teacher file lock guards cache metadata updates so that concurrent dataloader workers do not corrupt the cache state.

Table B.2: Logit-streaming settings.

Setting	Value
Teacher-logit storage	Cloudflare R2 object storage
Teacher-logit format	Raw per-sample <code>.pt</code> files
Remote path template	<code>remote_root/</code> <code><dataset>_teacher_logits_</code> <code><alias>/<teacher-slug>/</code> <code><index>.pt</code> <code>/tmp/teacher-logits-</code> <code>streaming</code> for R2 reads
Local cache directory	<code>/workspace/cache</code> for local-cache reads
Local cache limit	100 GB